

optimization_marketing_campaigns

June 6, 2025

Marketing Campaign Optimization

Project Goal

The project aims to optimize marketing campaigns by analyzing key factors affecting conversion, identifying the best promotion strategies, and building an effective model to predict the success of marketing activities.

Project Stages

1. Data Preparation

Removing duplicates Handling outliers and missing values Examining variable distributions

2. Exploratory Analysis

Identifying the most successful strategies that lead to high conversion rates. Determining key influencing factors: advertising channels, demographics, engagement. Analyzing correlations between variables to understand dependencies.

3. Budget Optimization

Budget allocation across different campaign types and channels using weights and factor importance for conversion

Features

Demographic Information

CustomerID: Unique identifier for each customer. Age: Age of the customer. Gender: Gender of the customer (Male/Female). Income: Annual income of the customer in USD.

Marketing-specific Variables

CampaignChannel: The channel through which the marketing campaign is delivered (Email, Social Media, SEO, PPC, Referral). CampaignType: Type of the marketing campaign (Awareness, Consideration, Conversion, Retention). AdSpend: Amount spent on the marketing campaign in USD. ClickThroughRate: Rate at which customers click on the marketing content. ConversionRate: Rate at which clicks convert to desired actions (e.g., purchases). AdvertisingPlatform: Confidential. AdvertisingTool: Confidential.

Customer Engagement Variables

WebsiteVisits: Number of visits to the website. PagesPerVisit: Average number of pages visited per session. TimeOnSite: Average time spent on the website per visit (in minutes). SocialShares: Number of times the marketing content was shared on social media. EmailOpens: Number of times marketing emails were opened. EmailClicks: Number of times links in marketing emails were clicked.

Historical Data

PreviousPurchases: Number of previous purchases made by the customer. LoyaltyPoints: Number of loyalty points accumulated by the customer.

Target Variable

Conversion: Binary variable indicating whether the customer converted (1) or not (0).

```
[538]: import numpy as np
import pandas as pd
import scipy as sp
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
from sklearn.preprocessing import RobustScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split, cross_val_score, \
    GridSearchCV
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix, precision_recall_curve, \
    roc_auc_score, roc_curve, fbeta_score, cohen_kappa_score, matthews_corrcoef, auc
from xgboost import XGBClassifier
from sklearn.ensemble import RandomForestClassifier
```

```
[540]: df = pd.read_csv('digital_marketing_campaign_dataset.csv')
```

```
[542]: df
```

```
[542]:
```

	CustomerID	Age	Gender	Income	CampaignChannel	CampaignType	\
0	8000	56	Female	136912	Social Media	Awareness	
1	8001	69	Male	41760	Email	Retention	
2	8002	46	Female	88456	PPC	Awareness	
3	8003	32	Female	44085	PPC	Conversion	
4	8004	60	Female	83964	PPC	Conversion	
...	
7995	15995	21	Male	24849	Email	Awareness	
7996	15996	43	Female	44718	SEO	Retention	
7997	15997	28	Female	125471	Referral	Consideration	
7998	15998	19	Female	107862	PPC	Consideration	
7999	15999	31	Female	93002	Email	Awareness	

	AdSpend	ClickThroughRate	ConversionRate	WebsiteVisits	\
0	6497.870068	0.043919	0.088031	0	
1	3898.668606	0.155725	0.182725	42	
2	1546.429596	0.277490	0.076423	2	
3	539.525936	0.137611	0.088004	47	
4	1678.043573	0.252851	0.109940	0	
...	

7995	8518.308575	0.243792	0.116773	23
7996	1424.613446	0.236740	0.190061	49
7997	4609.534635	0.056526	0.133826	35
7998	9476.106354	0.023961	0.138386	49
7999	7743.627070	0.185670	0.057228	15

	PagesPerVisit	TimeOnSite	SocialShares	EmailOpens	EmailClicks	\
0	2.399017	7.396803	19	6	9	
1	2.917138	5.352549	5	2	7	
2	8.223619	13.794901	0	11	2	
3	4.540939	14.688363	89	2	2	
4	2.046847	13.993370	6	6	6	
...	
7995	9.693602	14.227794	70	13	6	
7996	9.499010	3.501106	52	13	1	
7997	2.853241	14.618323	38	16	0	
7998	1.002964	3.876623	86	1	5	
7999	6.964739	12.763660	2	18	9	

	PreviousPurchases	LoyaltyPoints	AdvertisingPlatform	AdvertisingTool	\
0	4	688	IsConfid	ToolConfid	
1	2	3459	IsConfid	ToolConfid	
2	8	2337	IsConfid	ToolConfid	
3	0	2463	IsConfid	ToolConfid	
4	8	4345	IsConfid	ToolConfid	
...	
7995	7	286	IsConfid	ToolConfid	
7996	5	1502	IsConfid	ToolConfid	
7997	3	738	IsConfid	ToolConfid	
7998	7	2709	IsConfid	ToolConfid	
7999	9	341	IsConfid	ToolConfid	

	Conversion
0	1
1	1
2	1
3	1
4	1
...	...
7995	0
7996	0
7997	1
7998	1
7999	0

[8000 rows x 20 columns]

Data preparation

```
[545]: df = df.drop(columns=["AdvertisingPlatform", "AdvertisingTool"])
```

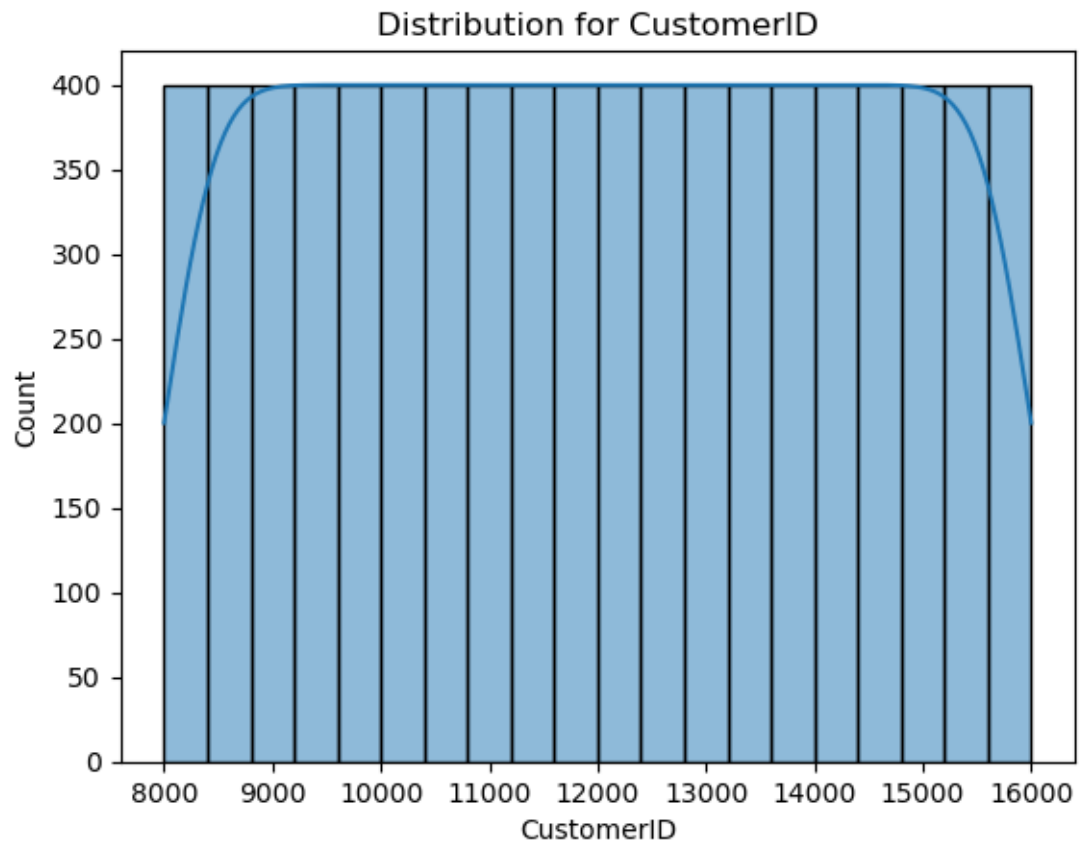
```
[547]: df.isnull().sum()
```


```
[547]: CustomerID      0
      Age            0
      Gender         0
      Income         0
      CampaignChannel 0
      CampaignType    0
      AdSpend         0
      ClickThroughRate 0
      ConversionRate  0
      WebsiteVisits   0
      PagesPerVisit   0
      TimeOnSite      0
      SocialShares    0
      EmailOpens       0
      EmailClicks      0
      PreviousPurchases 0
      LoyaltyPoints    0
      Conversion       0
      dtype: int64
```


```
[549]: df.duplicated().sum()
```


```
[549]: 0
```


```
[551]: for column in df.columns:
      sns.histplot(df[column].dropna(), kde=True)
      plt.title(f'Distribution for {column}')
      plt.show()
```




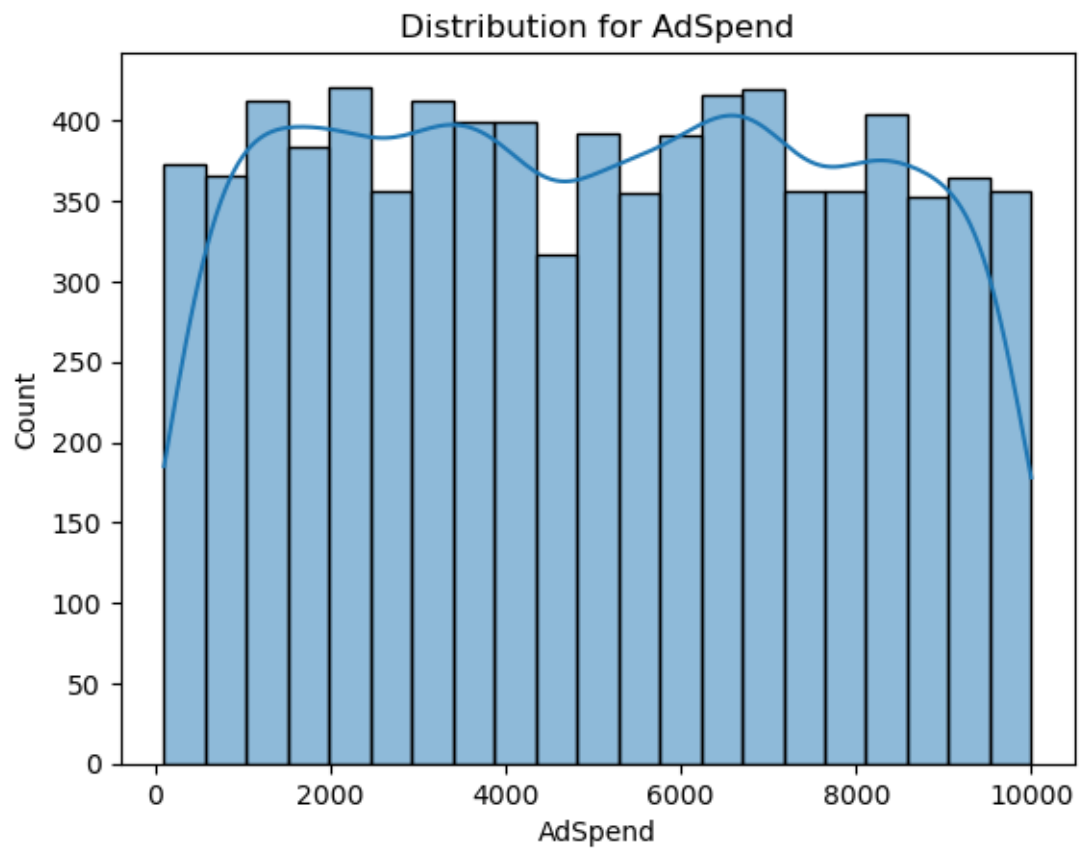


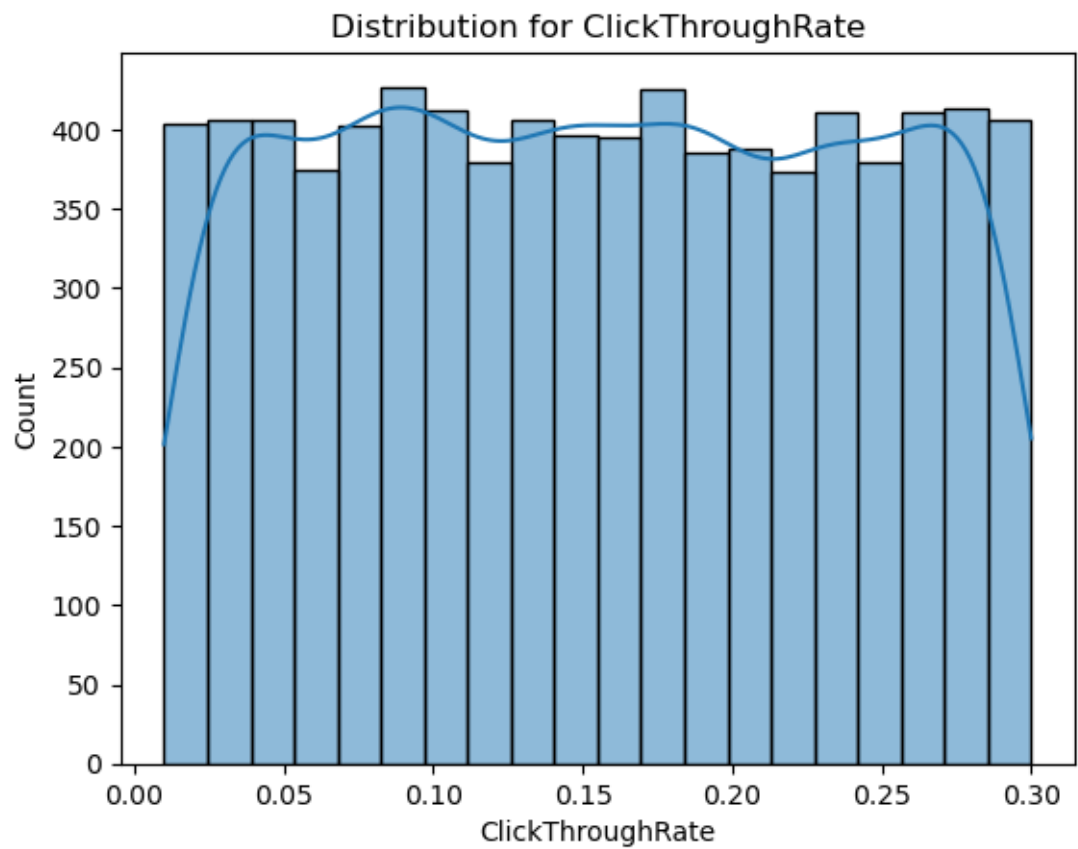


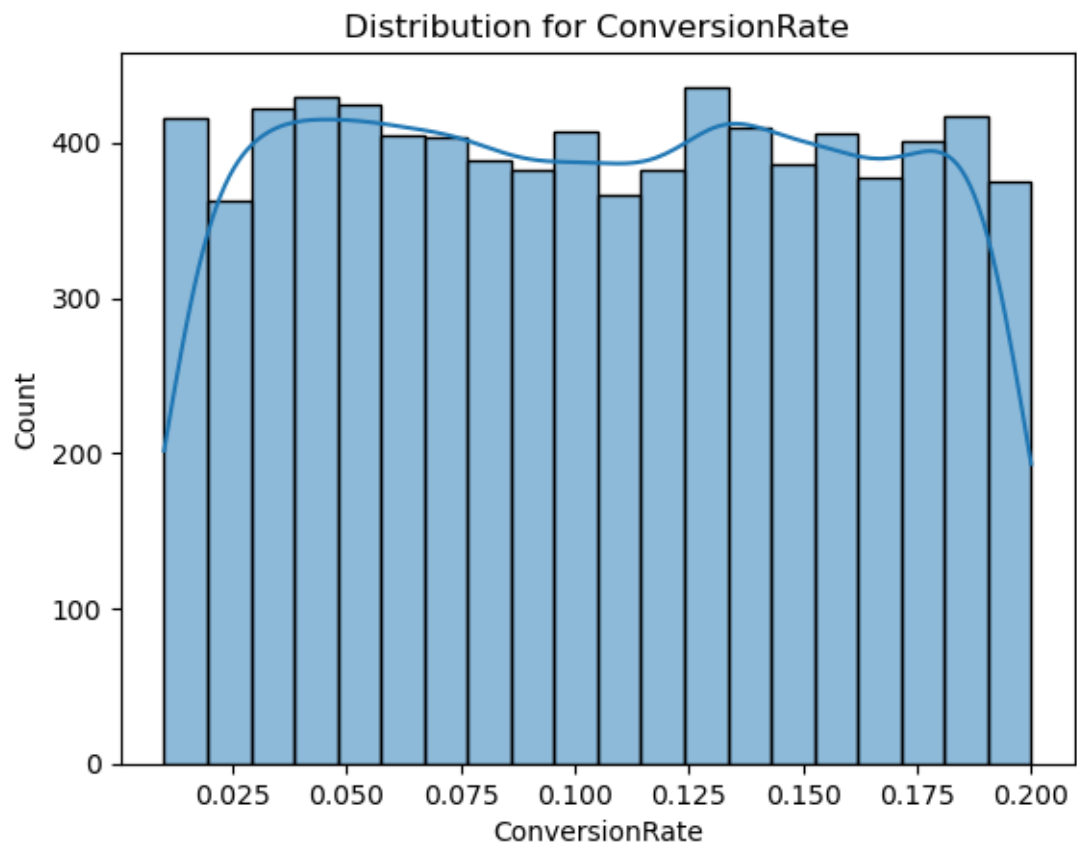


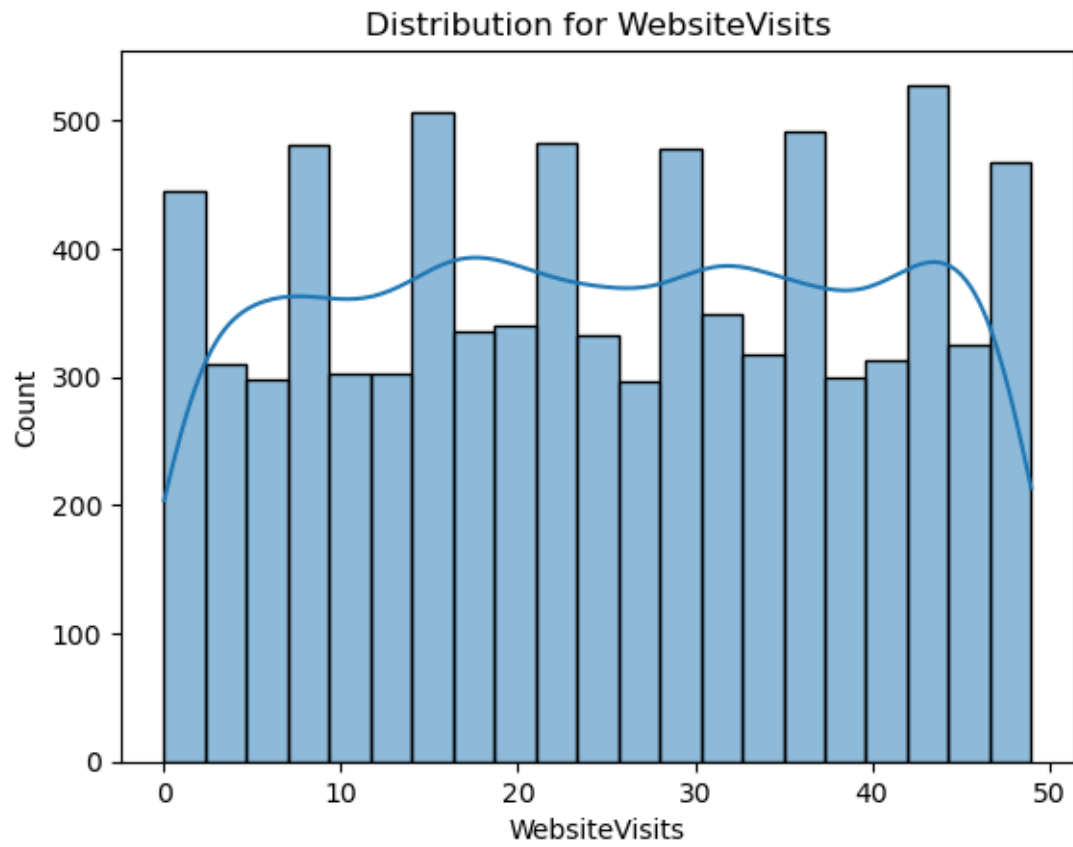


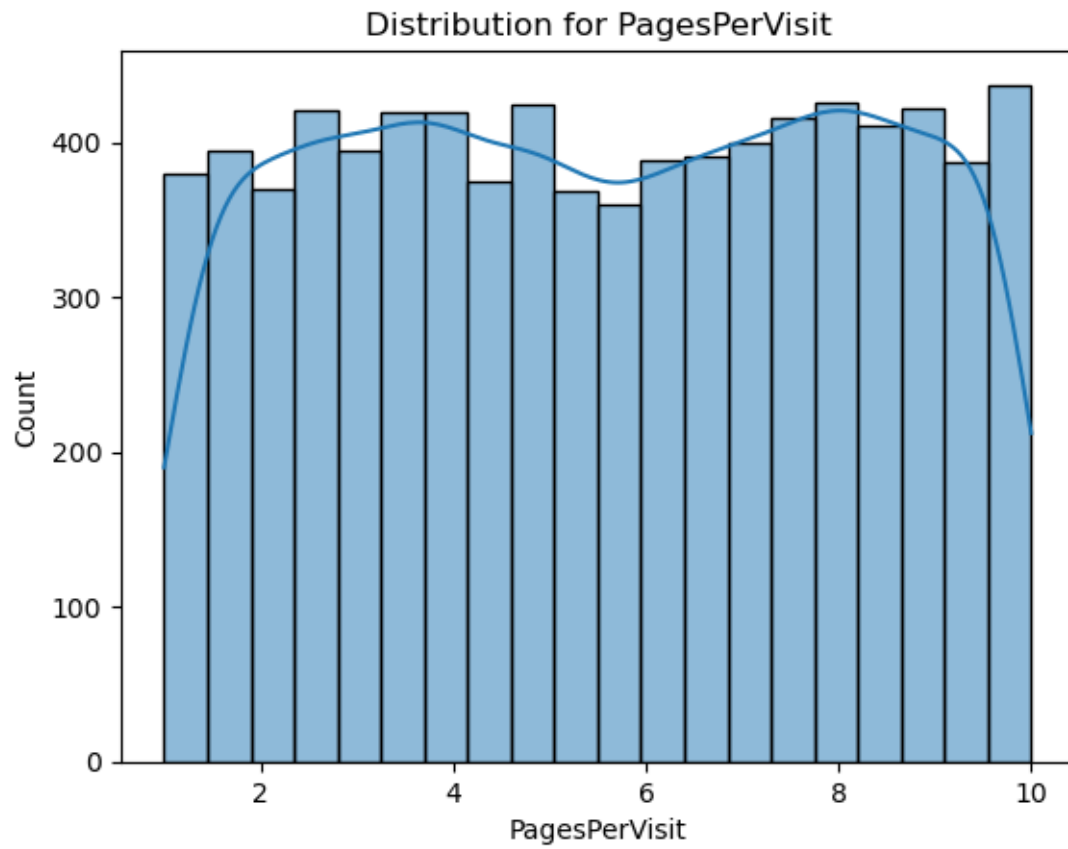


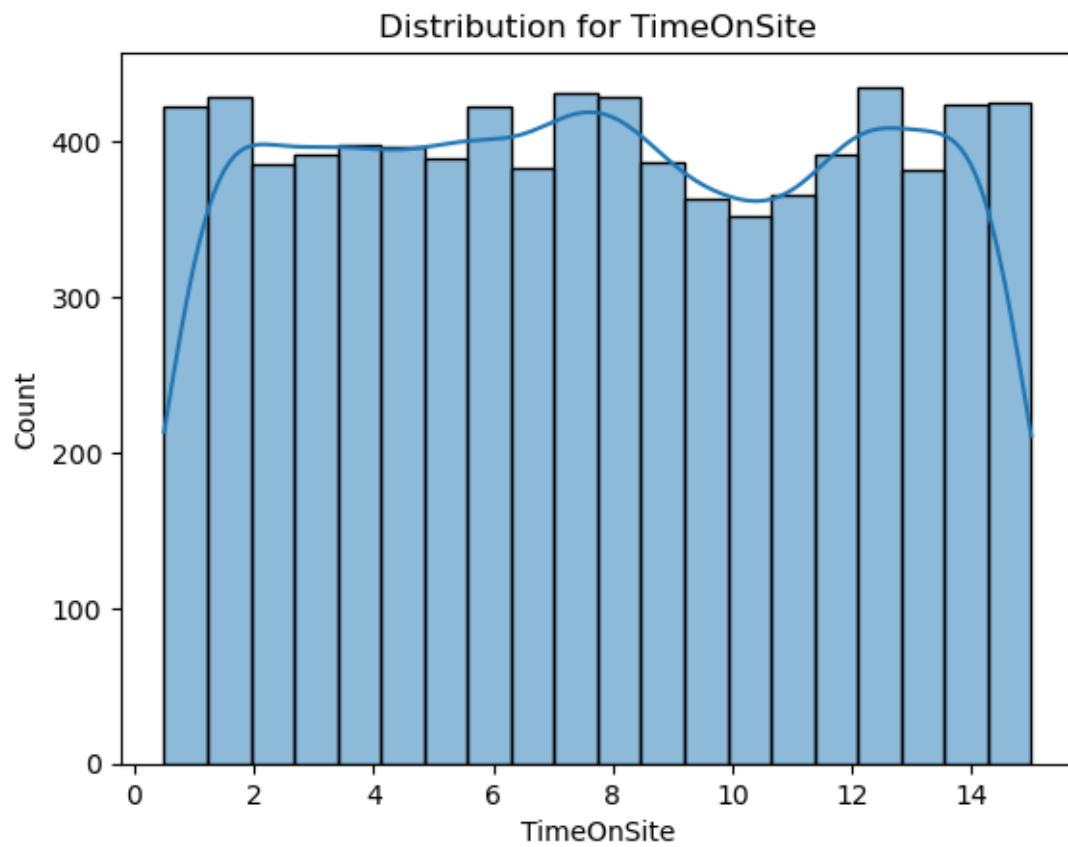


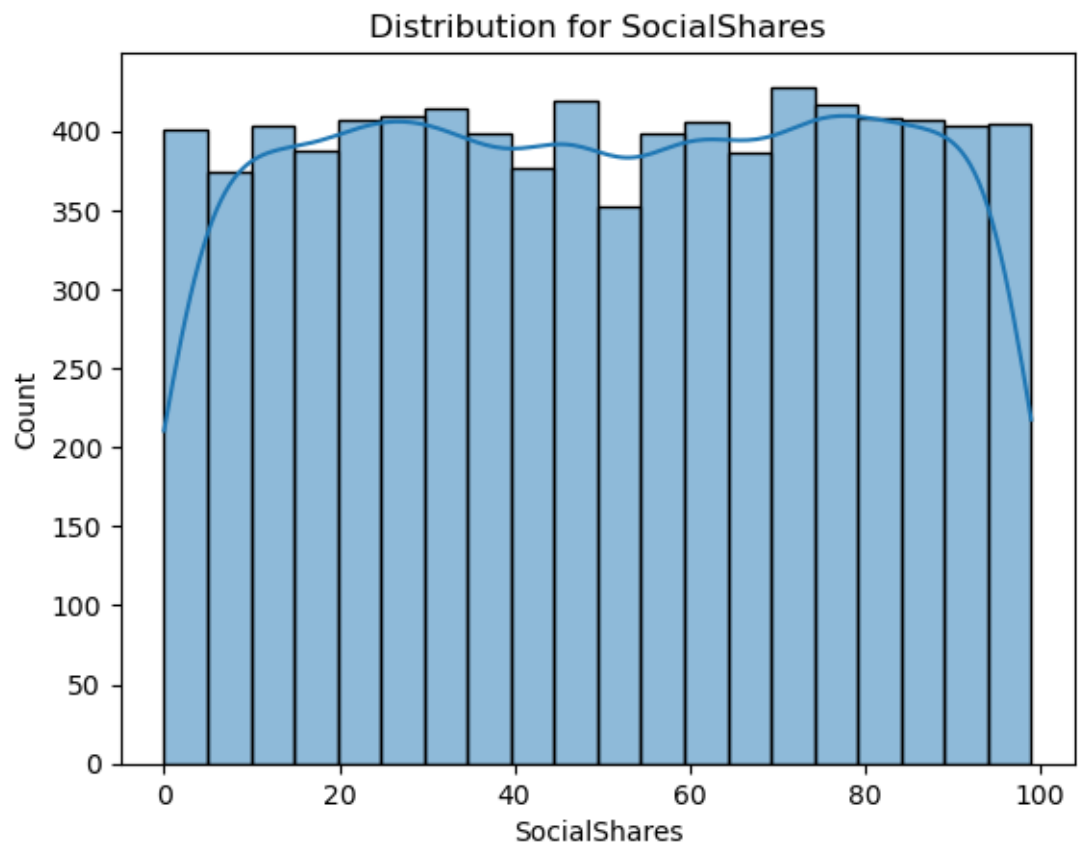


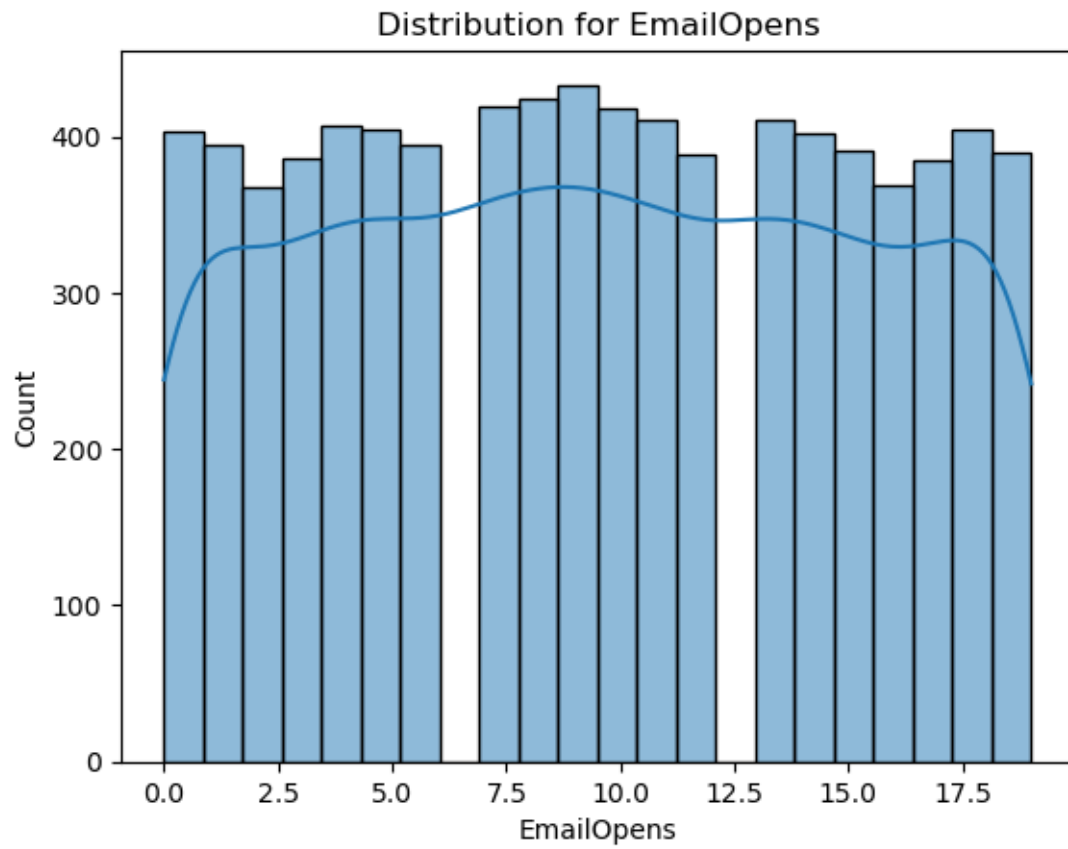


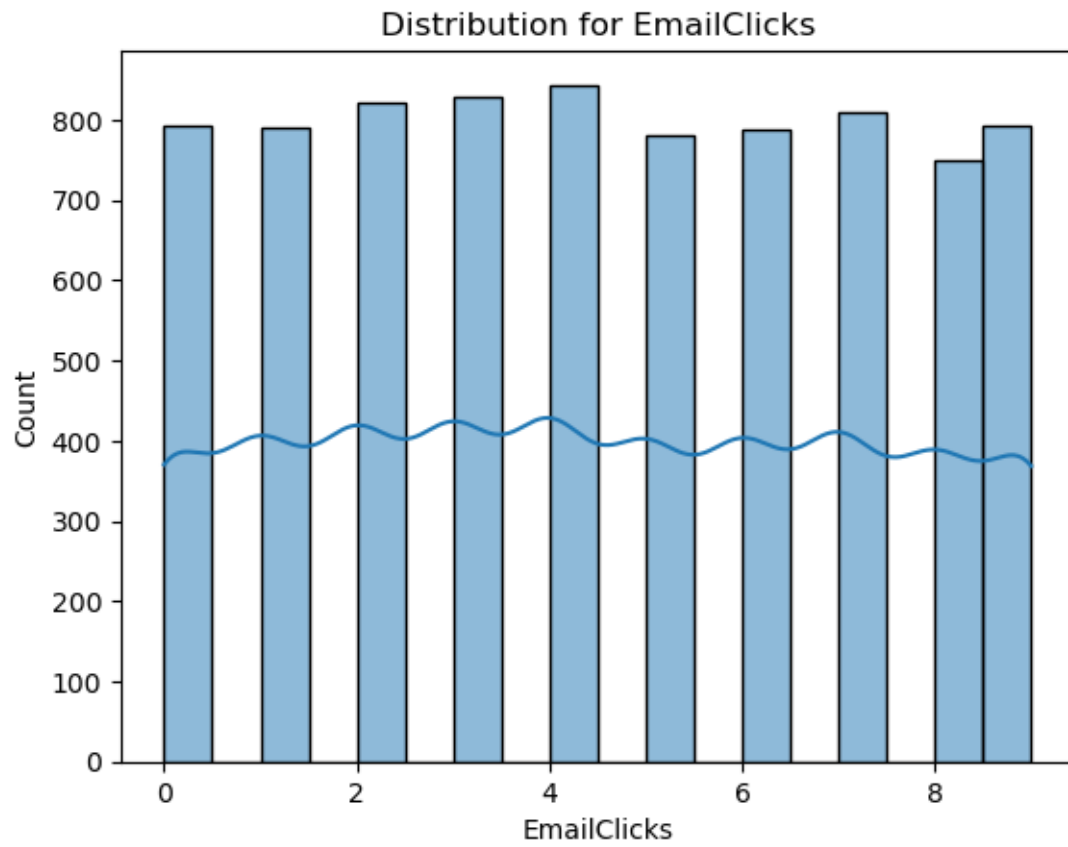



















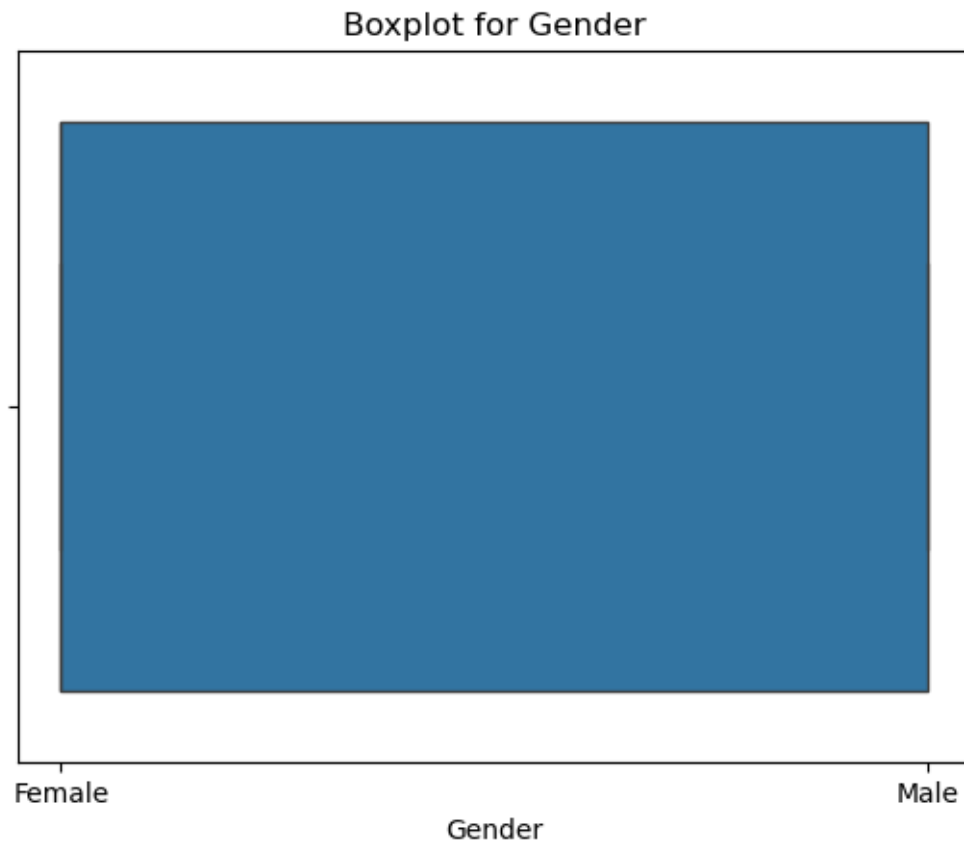





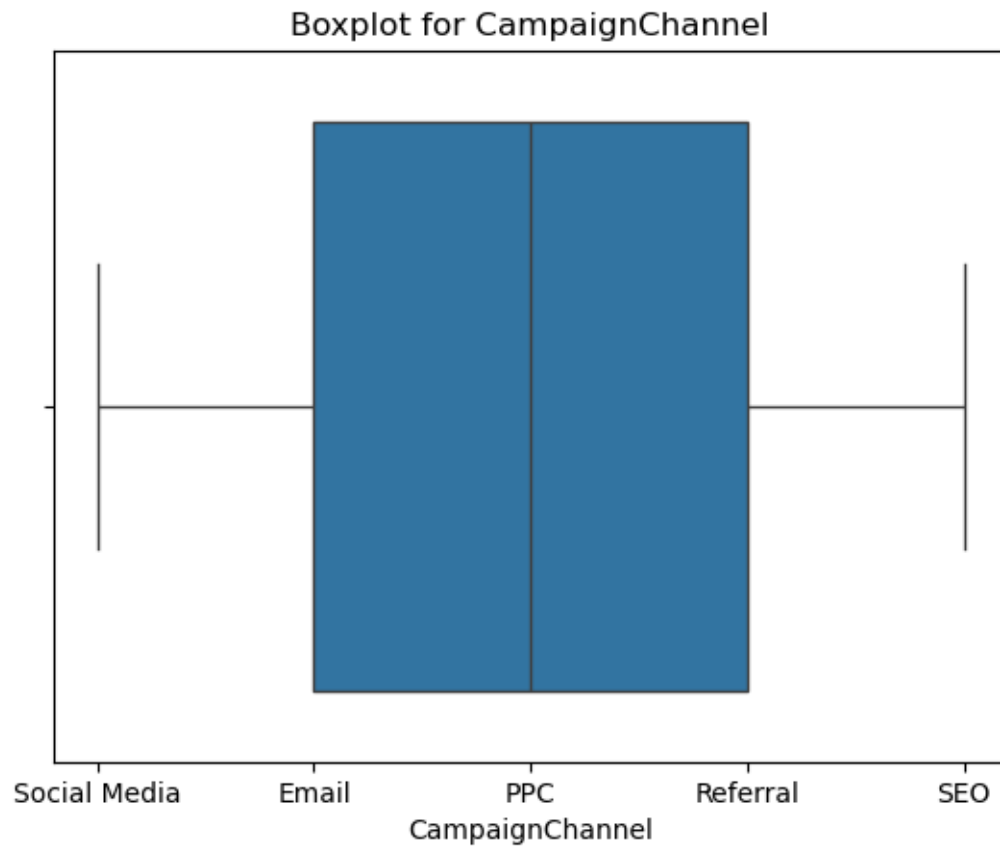
```
[552]: for column in df.columns:  
        sns.boxplot(data=df, x=column)  
        plt.title(f'Boxplot for {column}')  
        plt.show()
```

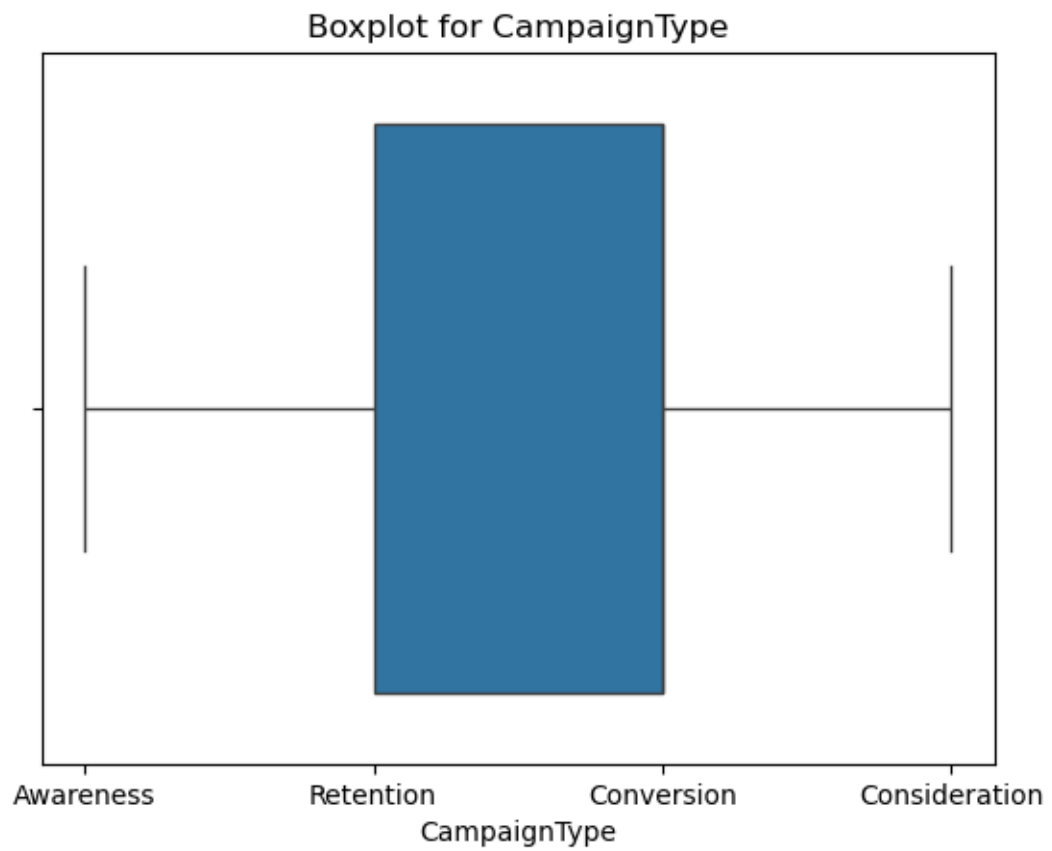


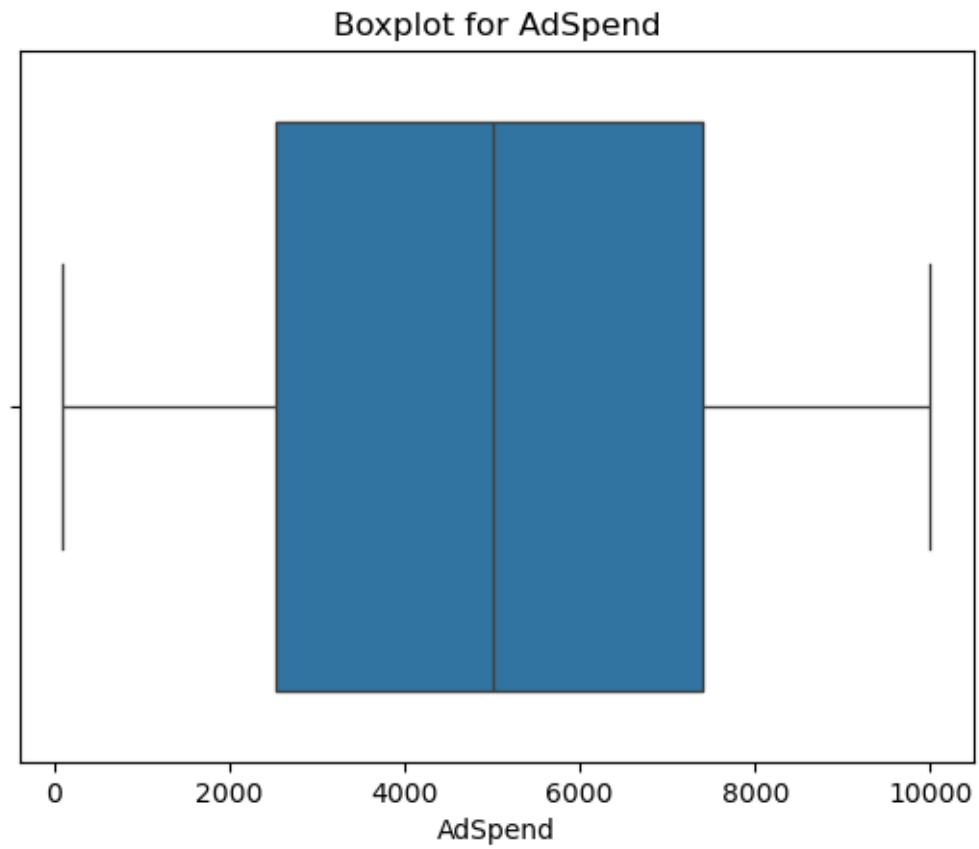





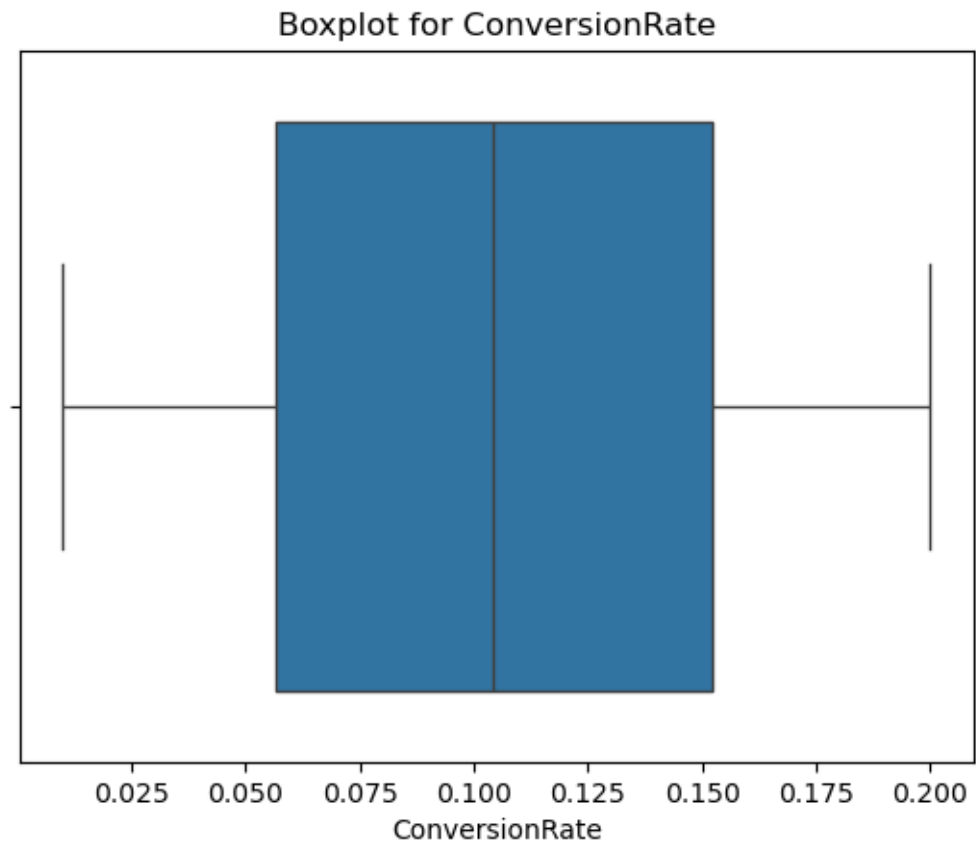


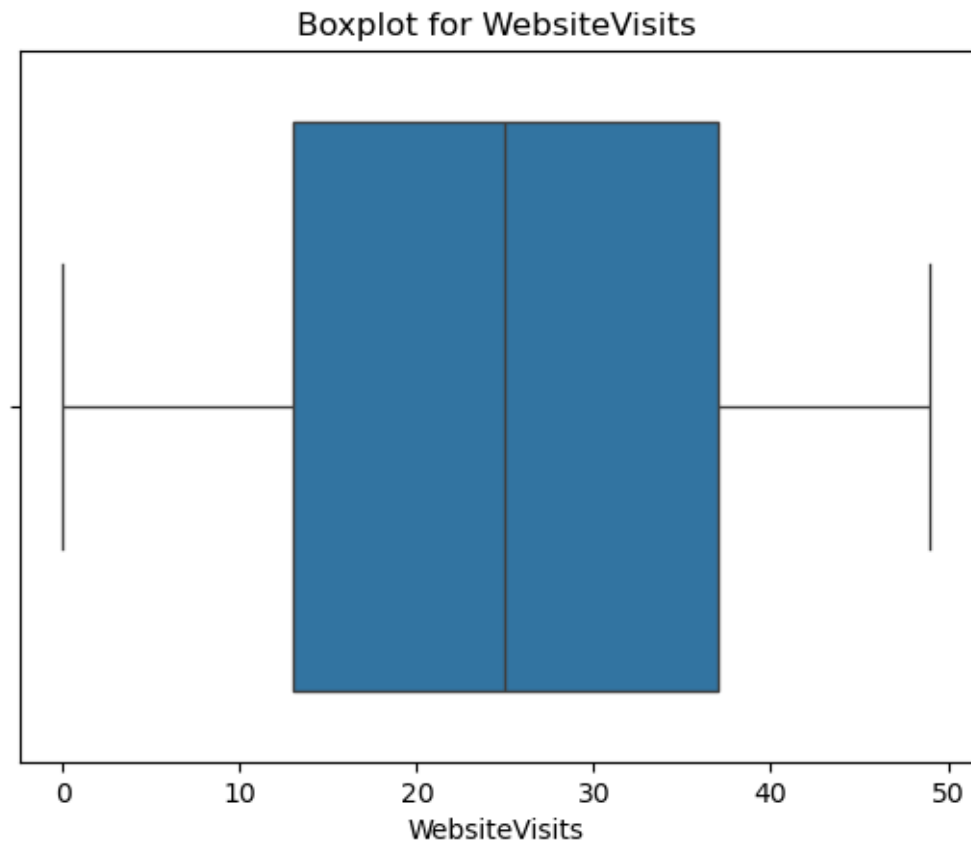


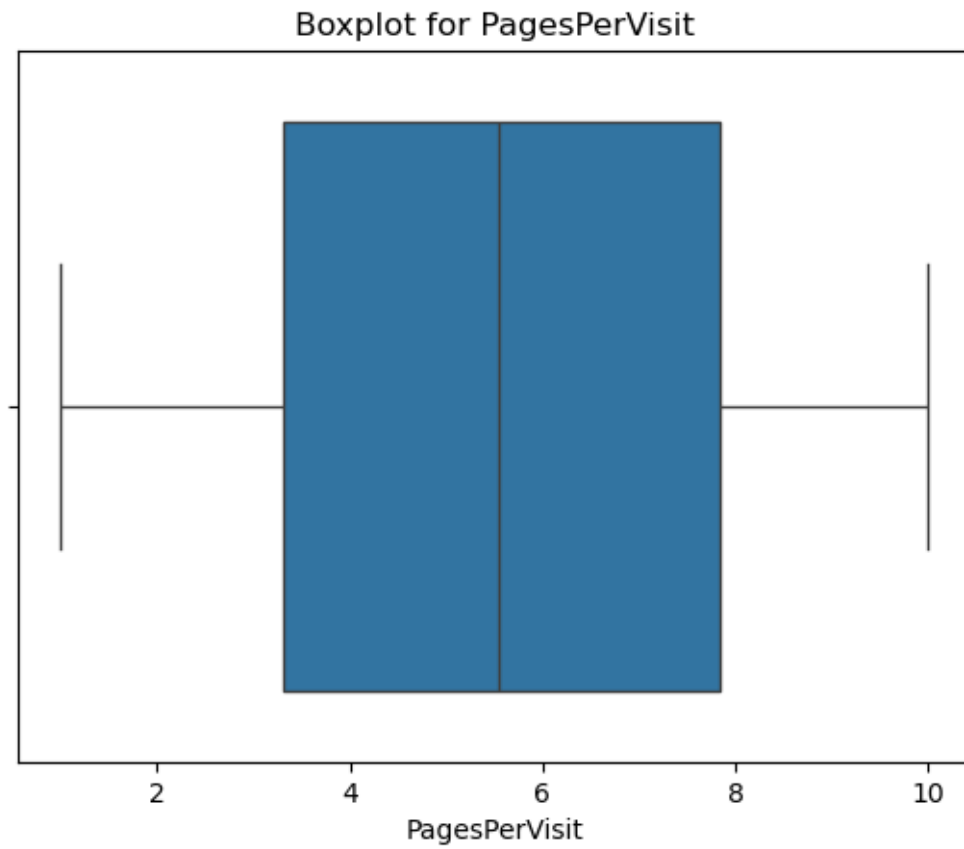


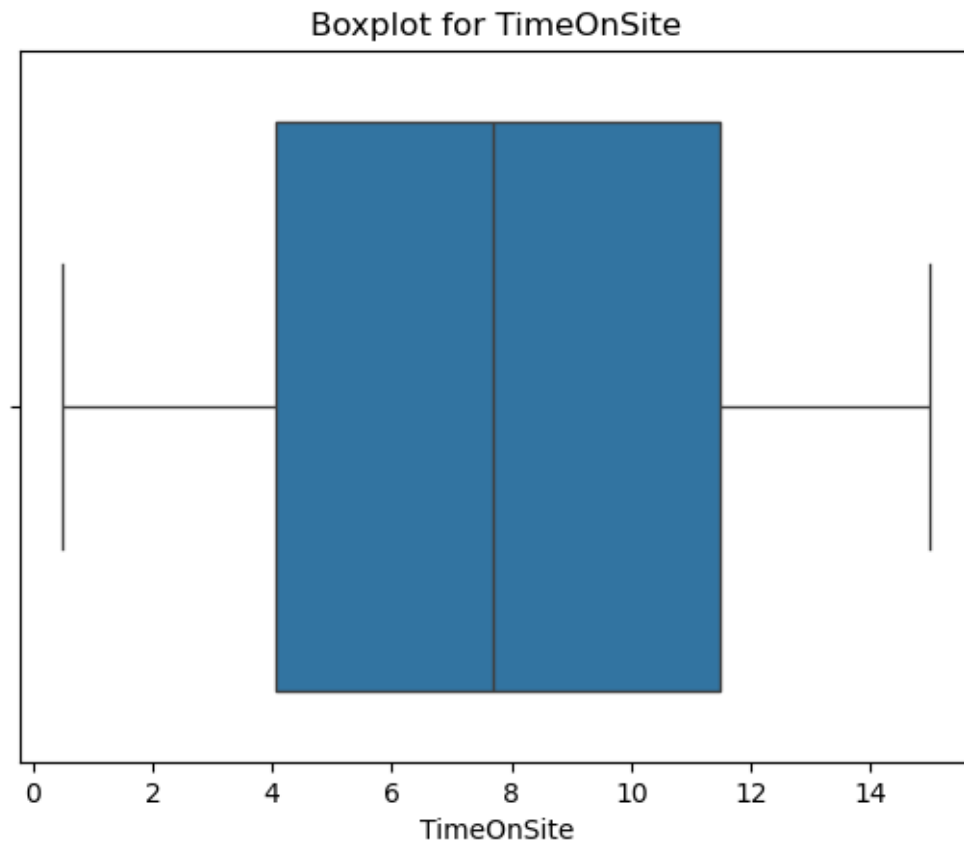


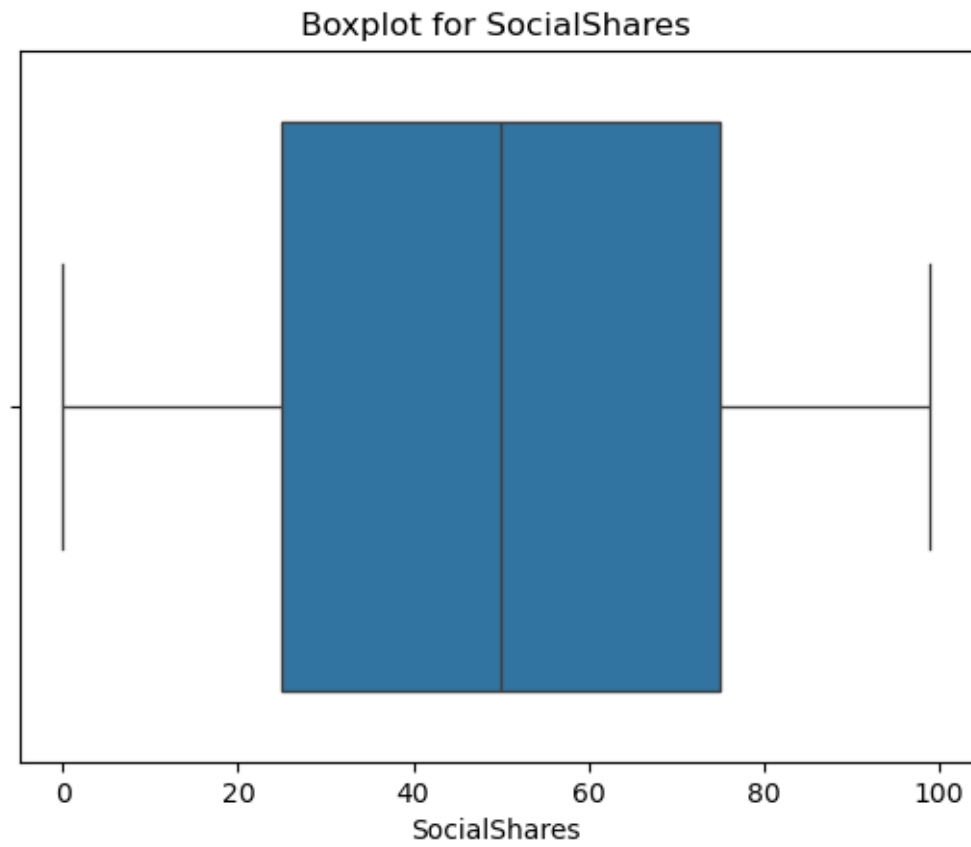


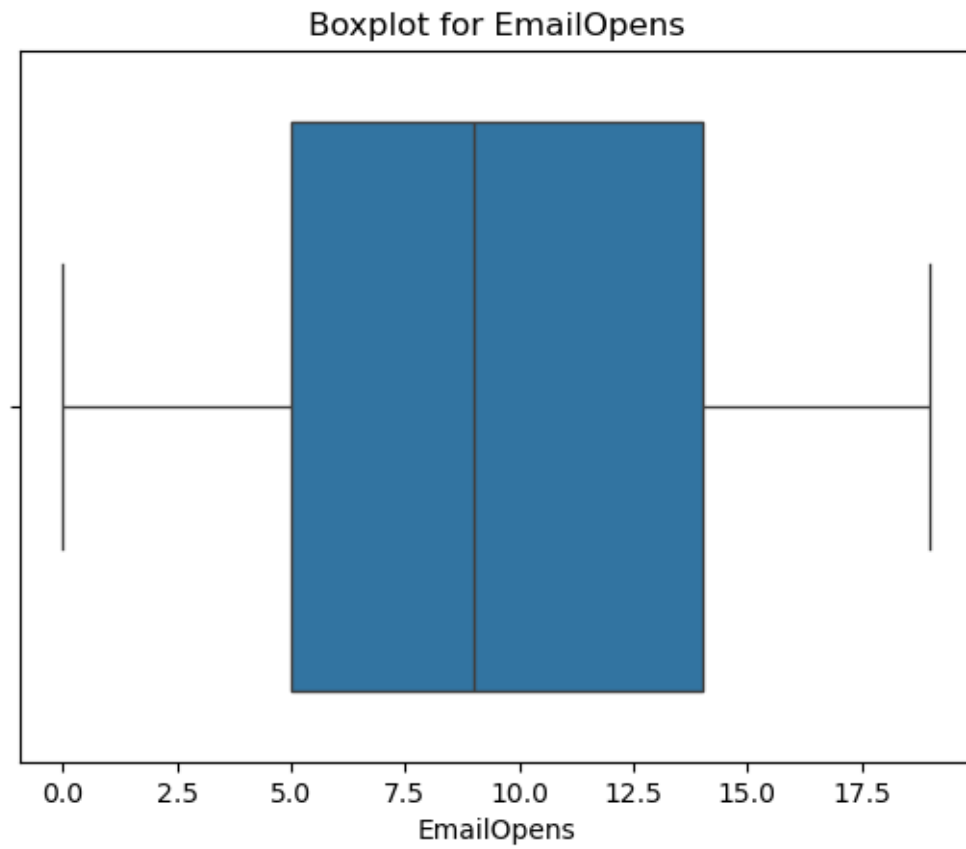
























Exploratory analysis

Questions to answer:

Correlations and relationships between variables

How are advertising expenses (AdSpend) related to conversion (Conversion)? How does customer engagement (WebsiteVisits, TimeOnSite) change with increased advertising expenses?

Segmentation and user grouping

Which customer groups are more likely to make purchases? Does campaign type (CampaignType) affect conversion across different age groups? What is the average spending (AdSpend) for each customer segment?

Email marketing effectiveness

Which age groups open email campaigns (EmailOpens) more frequently? Does the type of marketing campaign (CampaignType) influence email click rates (EmailClicks)? How many email openings (EmailOpens) lead to conversion?

Evaluation of marketing strategy effectiveness

Which marketing channels (CampaignChannel) demonstrate the highest efficiency? Which campaigns attract new customers, and which are more effective for retention?

Conversion funnel

Which advertising campaigns (CampaignType) drive the most traffic to the website? Which marketing channels (CampaignChannel) generate the highest website traffic? What percentage of visitors (WebsiteVisits) view more than one page (PagesPerVisit)? How does time spent on the site (TimeOnSite) affect the likelihood of conversion? Which advertising campaigns (CampaignType) contribute to longer site visits? Which marketing channels (CampaignChannel) encourage users to spend more time on the website? Which advertising campaigns (CampaignType) lead to more page views? Which marketing channels (CampaignChannel) result in higher page views? How does the number of website visits affect conversion? Visualization of the conversion funnel Which factors are most important for conversion?

```
[555]: df.describe()
```

```
[555]:
```

	CustomerID	Age	Income	AdSpend	ClickThroughRate \
count	8000.00000	8000.000000	8000.000000	8000.000000	8000.000000
mean	11999.50000	43.625500	84664.196750	5000.944830	0.154829
std	2309.54541	14.902785	37580.387945	2838.038153	0.084007
min	8000.00000	18.000000	20014.000000	100.054813	0.010005
25%	9999.75000	31.000000	51744.500000	2523.221165	0.082635
50%	11999.50000	43.000000	84926.500000	5013.440044	0.154505
75%	13999.25000	56.000000	116815.750000	7407.989369	0.228207
max	15999.00000	69.000000	149986.000000	9997.914781	0.299968

	ConversionRate	WebsiteVisits	PagesPerVisit	TimeOnSite \
count	8000.000000	8000.000000	8000.000000	8000.000000
mean	0.104389	24.751625	5.549299	7.727718
std	0.054878	14.312269	2.607358	4.228218
min	0.010018	0.000000	1.000428	0.501669
25%	0.056410	13.000000	3.302479	4.068340
50%	0.104046	25.000000	5.534257	7.682956
75%	0.152077	37.000000	7.835756	11.481468
max	0.199995	49.000000	9.999055	14.995311

	SocialShares	EmailOpens	EmailClicks	PreviousPurchases \
count	8000.000000	8000.000000	8000.000000	8000.000000
mean	49.799750	9.476875	4.467375	4.485500
std	28.901165	5.711111	2.856564	2.888093
min	0.000000	0.000000	0.000000	0.000000
25%	25.000000	5.000000	2.000000	2.000000
50%	50.000000	9.000000	4.000000	4.000000
75%	75.000000	14.000000	7.000000	7.000000
max	99.000000	19.000000	9.000000	9.000000

	LoyaltyPoints	Conversion
count	8000.000000	8000.000000
mean	2490.268500	0.876500
std	1429.527162	0.329031

```

min          0.000000    0.000000
25%         1254.750000    1.000000
50%         2497.000000    1.000000
75%         3702.250000    1.000000
max          4999.000000    1.000000

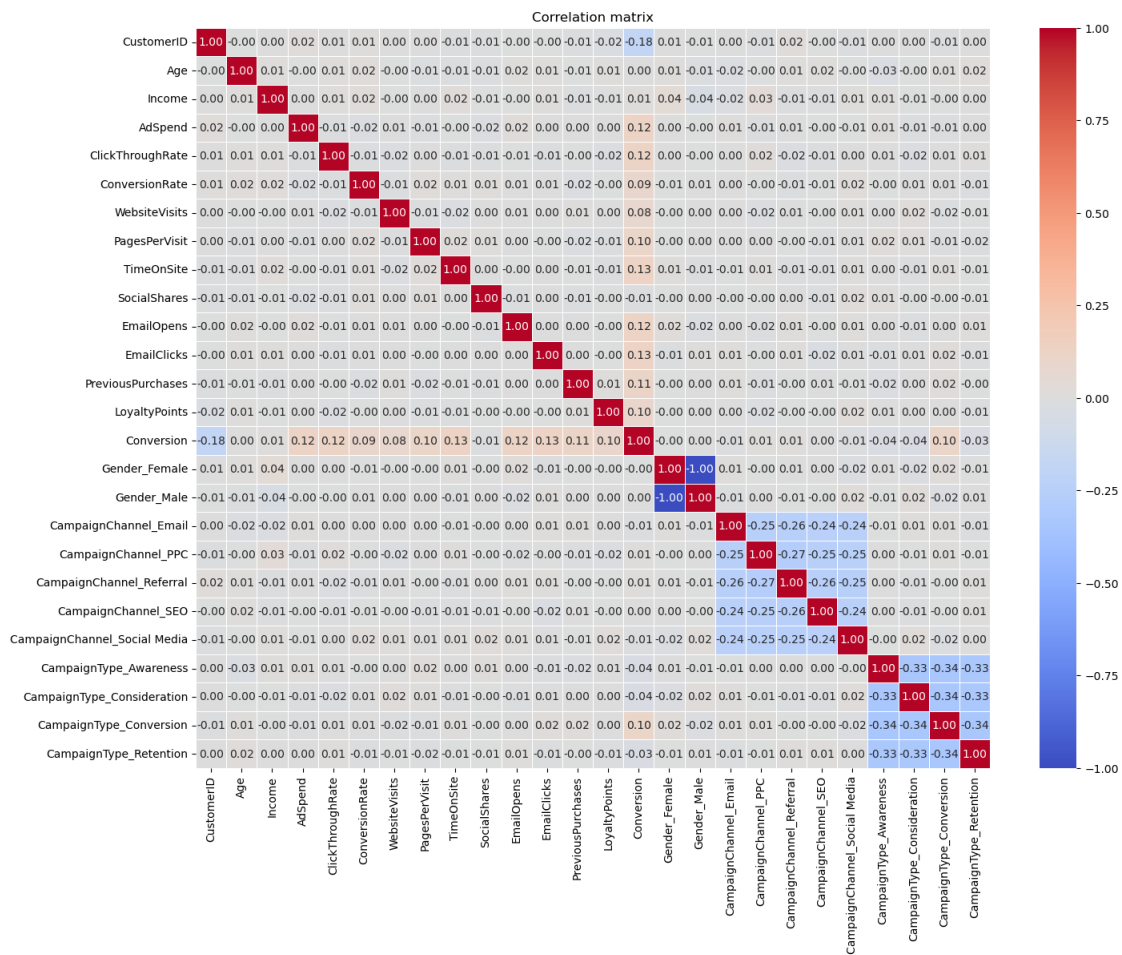
```

```

[556]: df_numeric = pd.get_dummies(df, drop_first=False)
corr_matrix = df_numeric.corr()

plt.figure(figsize=(16, 12))
sns.heatmap(corr_matrix, annot=True, fmt=".2f", cmap="coolwarm", linewidths=0.5)
plt.title("Correlation matrix")
plt.show()

```



How are advertising expenses (AdSpend) related to conversion (Conversion)?

```
[558]: bins = [100, 2000, 4000, 6000, 8000, 10000]
labels = ['100-2K', '2K-4K', '4K-6K', '6K-8K', '8K-10K']
df['AdSpendGroup'] = pd.cut(df['AdSpend'], bins=bins, labels=labels)


grouped_data = df.groupby('AdSpendGroup')['Conversion'].mean().reset_index()

plt.figure(figsize=(10, 6))
sns.barplot(x='AdSpendGroup', y='Conversion', data=grouped_data,
            palette='coolwarm')
plt.xlabel("Advertising expense groups ($)")
plt.ylabel("Average conversion rate (%)")
plt.title("Relationship between advertising expenses and conversion rate")
plt.ylim(0.7, 1)
plt.show()
```

```
/var/folders/g_/dlksrxdd3pz88bqsmz91cx540000gn/T/ipykernel_13099/4283026523.py:5
: FutureWarning: The default of observed=False is deprecated and will be changed
to True in a future version of pandas. Pass observed=False to retain current
behavior or observed=True to adopt the future default and silence this warning.
    grouped_data = df.groupby('AdSpendGroup')['Conversion'].mean().reset_index()
/var/folders/g_/dlksrxdd3pz88bqsmz91cx540000gn/T/ipykernel_13099/4283026523.py:8
: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='AdSpendGroup', y='Conversion', data=grouped_data,
            palette='coolwarm')
```



How does customer engagement (WebsiteVisits, TimeOnSite) change with increased advertising expenses?

```
[560]: grouped_data = df.groupby('AdSpendGroup')['WebsiteVisits'].mean().reset_index()
plt.figure(figsize=(10, 6))
ax = sns.barplot(x='AdSpendGroup', y='WebsiteVisits', data=grouped_data,
                palette=['#4CAF50', '#2196F3', '#FFC107', '#ff6f61', '#9C27B0'])
plt.xlabel("Advertising expense groups ($)")
plt.ylabel("Average number of website clicks")
plt.title("Relationship between advertising expenses and number of website
        clicks")
plt.ylim(20, 27)
for bar in ax.containers:
    ax.bar_label(bar, fmt='%.2f', label_type='edge', padding=3, color='black',
        fontsize=12)

plt.show()
```


```
/var/folders/g_/dlksrxdd3pz88bqsmz91cx540000gn/T/ipykernel_13099/356601795.py:1:
FutureWarning: The default of observed=False is deprecated and will be changed
to True in a future version of pandas. Pass observed=False to retain current
behavior or observed=True to adopt the future default and silence this warning.
```

```
grouped_data =
df.groupby('AdSpendGroup')['WebsiteVisits'].mean().reset_index()
/var/folders/g_/dlksrxdd3pz88bqsmz91cx540000gn/T/ipykernel_13099/356601795.py:3:
```

FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
ax = sns.barplot(x='AdSpendGroup', y='WebsiteVisits', data=grouped_data,
palette=['#4CAF50', '#2196F3', '#FFC107', '#ff6f61', '#9C27B0'])
```



```
[561]: grouped_data = df.groupby('AdSpendGroup')['TimeOnSite'].mean().reset_index()
plt.figure(figsize=(10, 6))
ax = sns.barplot(x='AdSpendGroup', y='TimeOnSite', data=grouped_data,
    palette=['#4CAF50', '#2196F3', '#FFC107', '#ff6f61', '#9C27B0'])
plt.xlabel("Advertising expense groups ($)")
plt.ylabel("Average time spent on the website (minutes)")
plt.title("Relationship between advertising expenses and time spent on the website")

plt.ylim(7, 8)
for bar in ax.containers:
    ax.bar_label(bar, fmt='%.2f', label_type='edge', padding=3, color='black',
        fontsize=12)

plt.show()
```

```

/var/folders/g_/dlksrxdd3pz88bqsmz91cx540000gn/T/ipykernel_13099/1096504721.py:1
: FutureWarning: The default of observed=False is deprecated and will be changed
to True in a future version of pandas. Pass observed=False to retain current
behavior or observed=True to adopt the future default and silence this warning.
grouped_data = df.groupby('AdSpendGroup')['TimeOnSite'].mean().reset_index()
/var/folders/g_/dlksrxdd3pz88bqsmz91cx540000gn/T/ipykernel_13099/1096504721.py:3
: FutureWarning:


```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

ax = sns.barplot(x='AdSpendGroup', y='TimeOnSite', data=grouped_data,
palette=['#4CAF50', '#2196F3', '#FFC107', '#ff6f61', '#9C27B0'])

```



Which customer groups are more likely to make purchases?

```

[563]: df['AgeGroup'] = pd.cut(df['Age'], bins=[18, 25, 35, 45, 55, 70],
↳ labels=['18-25', '26-35', '36-45', '46-55', '56+'])
df['IncomeGroup'] = pd.cut(df['Income'], bins=[20000, 50000, 80000, 110000,
↳ 150000], labels = ['Low', 'Medium', 'High', 'Very High'])
grouped_df = df.groupby(['Gender', 'AgeGroup',
↳ 'IncomeGroup'])['PreviousPurchases'].mean().reset_index()

```

```
top_10_groups = grouped_df.sort_values(by='PreviousPurchases', ascending=False).
↳head(10)
bottom_10_groups = grouped_df.sort_values(by='PreviousPurchases',
↳ascending=True).head(10)
```

/var/folders/g_/dlksrxdd3pz88bqsmz91cx540000gn/T/ipykernel_13099/270844897.py:3:
FutureWarning: The default of observed=False is deprecated and will be changed
to True in a future version of pandas. Pass observed=False to retain current
behavior or observed=True to adopt the future default and silence this warning.

```
grouped_df = df.groupby(['Gender', 'AgeGroup',
'IncomeGroup'])['PreviousPurchases'].mean().reset_index()
```

[564]: top_10_groups

```
[564]:      Gender AgeGroup IncomeGroup PreviousPurchases
14  Female    46-55         High         4.869767
16  Female     56+         Low         4.857143
10  Female    36-45         High         4.845494
28   Male    36-45         Low         4.785714
37   Male     56+        Medium         4.756906
29   Male    36-45        Medium         4.739130
4   Female    26-35         Low         4.725962
1   Female    18-25        Medium         4.709091
24   Male    26-35         Low         4.702532
23   Male    18-25      Very High         4.691057
```

[566]: bottom_10_groups

```
[566]:      Gender AgeGroup IncomeGroup PreviousPurchases
0   Female    18-25         Low         3.972028
38   Male     56+         High         4.023697
15  Female    46-55      Very High         4.077739
33   Male    46-55        Medium         4.240260
13  Female    46-55        Medium         4.271845
18  Female     56+         High         4.298361
39   Male     56+      Very High         4.300429
7   Female    26-35      Very High         4.305556
22   Male    18-25         High         4.316832
9   Female    36-45        Medium         4.317391
```

```
[567]: grouped_df = df.groupby(['Gender', 'AgeGroup', 'IncomeGroup'])['Conversion'].
↳mean().reset_index()

top_10_groups = grouped_df.sort_values(by='Conversion', ascending=False).
↳head(10)
bottom_10_groups = grouped_df.sort_values(by='Conversion', ascending=True).
↳head(10)
```

```

/var/folders/g_/dlksrxdd3pz88bqsmz91cx540000gn/T/ipykernel_13099/3436979725.py:1
: FutureWarning: The default of observed=False is deprecated and will be changed
to True in a future version of pandas. Pass observed=False to retain current
behavior or observed=True to adopt the future default and silence this warning.
grouped_df = df.groupby(['Gender', 'AgeGroup',
'IncomeGroup'])['Conversion'].mean().reset_index()

```

```
[568]: top_10_groups
```

```

[568]:      Gender AgeGroup IncomeGroup  Conversion
28    Male    36-45         Low    0.935714
29    Male    36-45        Medium    0.934783
9     Female   36-45        Medium    0.913043
11   Female   36-45    Very High    0.906250
26    Male    26-35         High    0.905109
31    Male    36-45    Very High    0.902174
30    Male    36-45         High    0.895105
18   Female     56+         High    0.895082
37    Male     56+        Medium    0.895028
27    Male    26-35    Very High    0.890710

```

```
[569]: bottom_10_groups
```

```

[569]:      Gender AgeGroup IncomeGroup  Conversion
22    Male    18-25         High    0.792079
33    Male    46-55        Medium    0.818182
1     Female   18-25        Medium    0.842424
35    Male    46-55    Very High    0.843373
4     Female   26-35         Low    0.846154
38    Male     56+         High    0.848341
0     Female   18-25         Low    0.860140
32    Male    46-55         Low    0.860606
24    Male    26-35         Low    0.860759
7     Female   26-35    Very High    0.861111

```

```

[570]: pivot_table = df.pivot_table(values=['PreviousPurchases', 'Conversion'],
    ↪ index=['AgeGroup', 'IncomeGroup'], columns=['Gender'])

plt.figure(figsize=(12, 6))
sns.heatmap(pivot_table['PreviousPurchases'], annot=True, fmt='.2f',
    ↪ cmap='coolwarm', linewidths=0.5)
plt.title("Intersection of groups by purchase frequency")
plt.xlabel("Gender")
plt.ylabel("Age and income level")
plt.show()

plt.figure(figsize=(12, 6))


```




```
sns.heatmap(pivot_table['Conversion'], annot=True, fmt='.2f', cmap='coolwarm',
            linewidths=0.5)
plt.title("Intersection of groups by conversion")
plt.xlabel("Gender")
plt.ylabel("Age and income level")
plt.show()
```

/var/folders/g_/dlksrxdd3pz88bqsmz91cx540000gn/T/ipykernel_13099/2636020496.py:1
: FutureWarning: The default value of observed=False is deprecated and will
change to observed=True in a future version of pandas. Specify observed=False to
silence this warning and retain the current behavior

```
pivot_table = df.pivot_table(values=['PreviousPurchases', 'Conversion'],
index=['AgeGroup', 'IncomeGroup'], columns=['Gender'])
```





1. Conversion by age and income

The highest conversion rate is observed in the 36-45 age group, especially among men with low and medium income levels (0.935 and 0.934). Women aged 36-45 with high income also show a high conversion rate (0.906). The lowest conversion rate is seen among men aged 18-25 with high income (0.792), which may indicate low engagement or longer decision-making times in this group.

2. Average number of previous purchases

Men aged 36-45 with low income make purchases most frequently (4.79). Women aged 56+ with low income (4.85) also exhibit high purchasing activity, possibly due to stable product needs. The lowest average number of purchases is seen among men aged 56+ with high income (4.02), which may suggest fewer but larger transactions.

3. Gender differences in purchasing behavior

Men aged 36-45 with low income are the most active buyers. Women aged 56+ make purchases consistently, regardless of income level. Women aged 18-25 with medium income make more purchases (4.7) than men (4.65), which may indicate a preference for frequent purchases.

4. Income impact

Medium income often demonstrates a higher conversion rate than high income (e.g., among 36-45-year-olds). High income does not always lead to more frequent purchases—groups with high income tend to shop less often.

Business conclusions:

The primary buyers are men aged 36-45 with low income and women aged 56+ with low income. Marketing campaigns for high conversion rates should target men aged 36-45 and women aged 56+. Personalization strategies can help improve engagement and purchasing behavior among men aged 56+. Women aged 18-25 with medium income tend to shop frequently, making them a strong target for subscription-based services.

Most efficient groups (high purchases + high conversion)

men aged 36-45 with low income (0.935 conversion, 4.79 purchases) — highest conversion rate, frequent buyers. women aged 36-45 with medium income (0.913 conversion, 4.31 purchases) — high conversion, stable purchasing behavior. men aged 26-35 with high income (0.905 conversion, 4.35 purchases) — strong conversion rate, relatively active buyers. women aged 56+ with low income (0.867 conversion, 4.85 purchases) — frequent shoppers, despite lower income, strong conversion. women aged 36-45 with high income (0.906 conversion, 4.84 purchases) — stable and active buyers.

Least efficient groups (low purchases + low conversion)

men aged 18-25 with high income (0.792 conversion, 4.31 purchases) — low conversion, infrequent purchases. women aged 46-55 with high income (0.865 conversion, 4.87 purchases) — frequent purchases, but lower conversion rate. men aged 56+ with high income (0.848 conversion, 4.02 purchases) — lowest purchase frequency. women aged 18-25 with low income (0.860 conversion, 3.97 purchases) — low purchasing volume. men aged 46-55 with medium income (0.818 conversion, 4.24 purchases) — low engagement levels.

Business conclusions:

focus on men aged 36-45 with low income, as they lead in purchases and conversions. women aged 56+ with low income are stable buyers, making them a valuable audience for mass-market products. men aged 18-25 with high income may require adjusted marketing strategies to better engage them. men aged 56+ with high income could be targeted for larger one-time purchases.

Does the campaign type (CampaignType) affect conversion across different age groups?

```
[576]: grouped_df = df.groupby(['AgeGroup', 'CampaignType'])['Conversion'].mean().
        ↪reset_index()
        grouped_df.sort_values(by='Conversion', ascending=False)

        plt.figure(figsize=(14, 6))
        ax = sns.barplot(x='AgeGroup', y='Conversion', hue='CampaignType', data=df,
        ↪palette=['#4CAF50', '#2196F3', '#FFC107', '#ff6f61', '#9C27B0'])

        for bar in ax.containers:
            ax.bar_label(bar, fmt='%.3f', label_type='edge', padding=12, color='black',
            ↪fontsize=8)


        plt.xlabel("Age group")
        plt.ylabel("Conversion (%)")
        plt.title("Impact of campaign type on conversion across different age groups")
        plt.legend(title="Campaign type")
        plt.show()
```

```
/var/folders/g_/dlksrxdd3pz88bqsmz91cx540000gn/T/ipykernel_13099/1876055778.py:1
: FutureWarning: The default of observed=False is deprecated and will be changed
to True in a future version of pandas. Pass observed=False to retain current
behavior or observed=True to adopt the future default and silence this warning.
    grouped_df = df.groupby(['AgeGroup',
```

```

'CampaignType']][ 'Conversion'].mean().reset_index()
/var/folders/g_/dlksrxdd3pz88bqsmz91cx540000gn/T/ipykernel_13099/1876055778.py:5
: UserWarning: The palette list has more values (5) than needed (4), which may
not be intended.
ax = sns.barplot(x='AgeGroup', y='Conversion', hue='CampaignType', data=df,
palette=['#4CAF50', '#2196F3', '#FFC107', '#ff6f61', '#9C27B0'])

```



```

[577]: grouped_df = df.groupby(['AgeGroup', 'CampaignChannel'])['Conversion'].mean().
        ↪reset_index()
grouped_df.sort_values(by='Conversion', ascending=False)

plt.figure(figsize=(14, 6))
ax = sns.barplot(x='AgeGroup', y='Conversion', hue='CampaignChannel', data=df,
        ↪palette=['#4CAF50', '#2196F3', '#FFC107', '#ff6f61', '#9C27B0'])

for bar in ax.containers:
    ax.bar_label(bar, fmt='%.3f', label_type='edge', padding=15, color='black',
        ↪fontsize=8)

plt.xlabel("Age group")
plt.ylabel("Conversion (%)")
plt.title("Impact of campaign channel on conversion across different age_
        ↪groups")
plt.legend(title="Campaign channel")

plt.show()


```

```

/var/folders/g_/dlksrxdd3pz88bqsmz91cx540000gn/T/ipykernel_13099/3952508184.py:1
: FutureWarning: The default of observed=False is deprecated and will be changed
to True in a future version of pandas. Pass observed=False to retain current

```

```
behavior or observed=True to adopt the future default and silence this warning.
grouped_df = df.groupby(['AgeGroup',
'CampaignChannel'])['Conversion'].mean().reset_index()
```



Most effective combinations (high conversion):

PPC (Pay-Per-Click Advertising) in the 36-45 age group (0.921) This group shows the highest conversion rate among all channels. PPC performs well for this age category, likely due to their tendency for targeted purchases. Referral marketing in the 36-45 age group (0.910) A high conversion rate suggests strong trust in recommendations, especially among people in this age bracket. SEO (Search Engine Optimization) in the 36-45 age group (0.898) This indicates that organic search plays a key role in their purchasing decisions. PPC in the 56+ age group (0.890) Paid ads are also effective for older audiences, likely because they prefer direct messaging and are ready to take action. SEO in the 18-25 age group (0.889) Young consumers frequently rely on search engines when choosing products or services.

Least effective combinations (low conversion):

Social Media for the 18-25 age group (0.831) Despite high social media usage, conversion rates are lower than other channels. Young users may engage with content but take longer to make purchasing decisions. PPC for the 26-35 age group (0.843) The conversion rate is lower than in other age categories, possibly indicating a more rational approach to purchases. Email marketing in the 18-25 age group (0.844) Younger consumers are less responsive to email campaigns, as their attention is often focused on other channels. SEO for the 46-55 age group (0.848) Search is less effective at converting buyers in this age bracket, possibly because they rely more on recommendations. Referral marketing for the 18-25 age group (0.853) Recommendations are less influential among young consumers, likely because they prefer independent research and comparisons.

General conclusions:

Best channels: PPC and referral marketing for the 36-45 age group. SEO delivers good results for younger consumers (18-25 years old). Social media works well for engagement but does not lead to

high conversions. Email is effective for older demographics but weaker for younger segments.

Strategy optimization:

Strengthen PPC for 36-45 and 56+ years. Increase trust-based marketing (Referral) for middle-age segments. Reassess Social Media and Email marketing for younger consumers.

Most effective campaign types (high conversion):

Conversion campaigns yield the best results across all age groups. The highest conversion rate is observed in the 26-35 age group (0.951), indicating their readiness to make purchases. 36-45 years (0.946) also exhibit strong conversion rates, confirming their decisiveness. 18-25 years (0.935) show high conversion, likely due to their impulse-driven purchasing behavior. 56+ years (0.931) respond well to conversion-focused campaigns, demonstrating a willingness to act.

Less effective campaigns (low conversion):

Awareness campaigns show the lowest conversion rates, particularly in the 18-25 age group (0.838). This suggests that brand-awareness-focused efforts do not immediately lead to purchases. 26-35 years (0.831) also show low conversion, possibly due to insufficient motivation to take action. Consideration campaigns similarly yield lower conversion rates, especially for the 18-25 group (0.821). This may indicate that young consumers tend to explore options but delay decision-making.

Other insights:

Retention campaigns are most effective for 36-45 and 56+ years, confirming that loyal older customers respond well to retention-focused strategies. Consideration and awareness campaigns play an important role for 46-55 and 56+ years, possibly because they make more thoughtful decisions and require additional information.

Business conclusions:

Prioritize conversion campaigns, especially for the 26-35 and 36-45 age groups. Awareness campaigns are less effective for younger consumers—direct offers or promotions may work better. Retention strategies are crucial for consumers aged 36+, as loyalty campaigns drive engagement. Optimizing consideration campaigns for the 18-25 age group through interactive and personalized offers could improve performance.

What is the average advertising spend (AdSpend) for each customer segment?

```
[590]: avg_ad_spend = df.groupby(['Gender', 'AgeGroup', 'IncomeGroup'])['AdSpend'].  
        ↪mean().reset_index()  
avg_ad_spend = avg_ad_spend.sort_values(by='AdSpend', ascending=False)  
avg_ad_spend.head(5)
```

```
/var/folders/g_/dlksrxdd3pz88bqsmz91cx540000gn/T/ipykernel_13099/2841502055.py:1  
: FutureWarning: The default of observed=False is deprecated and will be changed  
to True in a future version of pandas. Pass observed=False to retain current  
behavior or observed=True to adopt the future default and silence this warning.  
    avg_ad_spend = df.groupby(['Gender', 'AgeGroup',  
'IncomeGroup'])['AdSpend'].mean().reset_index()
```

```
[590]:
```

	Gender	AgeGroup	IncomeGroup	AdSpend
22	Male	18-25	High	5436.881631
4	Female	26-35	Low	5424.542116
20	Male	18-25	Low	5388.525293
24	Male	26-35	Low	5337.734579
9	Female	36-45	Medium	5299.330935

```
[593]: avg_ad_spend.tail(5)
```

```
[593]:
```

	Gender	AgeGroup	IncomeGroup	AdSpend
18	Female	56+	High	4779.197470
12	Female	46-55	Low	4745.416904
37	Male	56+	Medium	4741.535916
3	Female	18-25	Very High	4593.637646
1	Female	18-25	Medium	4438.523145

Most efficient groups (high conversion with reasonable expenses):

Men aged 36-45 (low and medium income) — 0.935 and 0.934 High conversion with moderate advertising expenses (not among the top spenders). These groups respond well to marketing campaigns, making them worth retaining. Women aged 36-45 (very high income) — 0.906 Excellent conversion rate, despite advertising expenses not being among the highest, making this group attractive from an investment perspective. Men aged 26-35 (high and very high income) — 0.905 and 0.890 A good balance between conversion and advertising expenses. Women aged 56+ (high income) — 0.895 High conversion but moderate expenses. This segment can yield stable profits with relatively low ad spending.

Least efficient groups (low conversion with high expenses):

Men aged 18-25 (high income) — 0.792, but AdSpend = 5436 High advertising expenses but lowest conversion rate among all groups. This strategy needs reconsideration, as it may be inefficient. Women aged 26-35 (low income) — 0.846, but AdSpend = 5424 Moderate conversion but very high advertising costs, making this segment less profitable. Men aged 18-25 (low income) — 0.860, but AdSpend = 5388 Similar to the previous case, above-average spending but unimpressive conversion. Women aged 36-45 (medium income) — 0.913, but AdSpend = 5299 Despite high conversion, advertising costs are also high. The strategy should be optimized to reduce expenses without losing effectiveness.

Which age groups open email campaigns (EmailOpens) more frequently?

```
[597]: bins = [18, 25, 35, 45, 55, 65, 100]
labels = ["18-24", "25-34", "35-44", "45-54", "55-64", "65+"]
df["AgeGroup"] = pd.cut(df["Age"], bins=bins, labels=labels, right=False)

email_opens_by_age = df.groupby("AgeGroup")["EmailOpens"].mean().
    ↪sort_values(ascending=False)

plt.figure(figsize=(12, 6))
ax = email_opens_by_age.plot(kind="bar", color="skyblue", edgecolor="black")
```

```

for bar in ax.containers:
    ax.bar_label(bar, fmt='%.3f', label_type='edge', padding=3, color='black',
    ↪ fontsize=12)


plt.xlabel("Age group")
plt.ylabel("Average number of email opens")
plt.title("Email open frequency by age group")
plt.xticks(rotation=0)
plt.ylim(9, 10)
plt.show()

```

```

/var/folders/g_/dlksrxdd3pz88bqsmz91cx540000gn/T/ipykernel_13099/3378274548.py:5
: FutureWarning: The default of observed=False is deprecated and will be changed
to True in a future version of pandas. Pass observed=False to retain current
behavior or observed=True to adopt the future default and silence this warning.
    email_opens_by_age =
df.groupby("AgeGroup")["EmailOpens"].mean().sort_values(ascending=False)

```



Does the type of marketing campaign (CampaignType) affect the email click-through rate (EmailClicks)?

```

[600]: email_clicks_by_campaign_type = df.groupby("CampaignType")["EmailClicks"].
    ↪ mean().sort_values(ascending=False)

plt.figure(figsize=(12, 6))
ax = email_clicks_by_campaign_type.plot(kind="bar", color="#4CAF50",
    ↪ edgecolor="black")

```




```

for bar in ax.containers:
    ax.bar_label(bar, fmt='%.3f', label_type='edge', padding=3, color='black',
    ↪fontsize=12)

plt.xlabel("Type of marketing campaign")
plt.ylabel("Average number of email link clicks")
plt.title("Impact of marketing campaign type on email link clicks")
plt.xticks(rotation=0)
plt.ylim(4, 5)
plt.show()

```



How many email opens (EmailOpens) lead to conversions?

```

[605]: converted_df = df[df["Conversion"] == 1]
average_email_opens = converted_df["EmailOpens"].mean()
average_email_opens

```

[605]: 9.744580718767827

Which marketing channels (CampaignChannel) demonstrate the best effectiveness?

```

[609]: df['Efficiency'] = df['Conversion'] / df['AdSpend']
grouped_df = df.groupby(['CampaignChannel'])['Efficiency'].mean().reset_index()
grouped_df_sorted = grouped_df.sort_values(by='Efficiency', ascending=False)

plt.figure(figsize=(12, 6))
ax = sns.barplot(x='CampaignChannel', y='Efficiency', data=grouped_df_sorted,
    ↪palette='coolwarm', order=grouped_df_sorted['CampaignChannel'])

```

```

for bar in ax.containers:
    ax.bar_label(bar, fmt='%.6f', label_type='edge', padding=3, color='black',
    ↪fontsize=12)

plt.xlabel("Marketing channel")
plt.ylabel("Average efficiency (conversion per $ spent)")
plt.title("Channel effectiveness (descending order)")
plt.xticks(rotation=45)
plt.ylim(0.0003, 0.0004)
plt.show()

```


/var/folders/g_/dlksrxdd3pz88bqsmz91cx540000gn/T/ipykernel_13099/1609171232.py:6
: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

ax = sns.barplot(x='CampaignChannel', y='Efficiency', data=grouped_df_sorted,
palette='coolwarm', order=grouped_df_sorted['CampaignChannel'])

```



Which campaigns attract new customers, and which are better for retention?

```

[612]: df["new_customer_ratio"] = df["ClickThroughRate"] * df["ConversionRate"]
df["retention_ratio"] = df["PreviousPurchases"] + df["LoyaltyPoints"]

```

```

new_customer_threshold = df["new_customer_ratio"].median()
retention_threshold = df["retention_ratio"].median()

result_df = df.groupby("CampaignType")[["new_customer_ratio",
↪ "retention_ratio"]].mean().reset_index()
result_df

```

```

[612]:
   CampaignType  new_customer_ratio  retention_ratio
0    Awareness           0.016103      2507.114185
1  Consideration           0.015974      2502.758048
2    Conversion           0.016298      2495.562831
3     Retention           0.016117      2473.098100

```

Thus, Conversion campaigns are ideal for actively attracting new customers. Awareness and Consideration campaigns help gradually engage customers and maintain their long-term interest in the brand.

Which advertising campaigns (CampaignType) generate the most website traffic?

```


[616]: traffic_by_campaign_type = df.groupby("CampaignType")["WebsiteVisits"].mean().
↪ sort_values(ascending=False)

plt.figure(figsize=(12, 6))
ax = traffic_by_campaign_type.plot(kind="bar", color='#9C27B0',
↪ edgecolor="black")

for bar in ax.containers:
    ax.bar_label(bar, fmt='%.3f', label_type='edge', padding=3, color='black',
↪ fontsize=12)

plt.xlabel("Type of advertising campaign")
plt.ylabel("Average number of website visits")
plt.title("Impact of advertising campaign type on website traffic")
plt.xticks(rotation=0)
plt.ylim(20, 30)
plt.show()

```




Which advertising channels (CampaignChannel) drive the most website traffic?

```
[619]: traffic_by_campaign_channel = df.groupby("CampaignChannel")["WebsiteVisits"].
        ↪mean().sort_values(ascending=False)

plt.figure(figsize=(12, 6))
ax = traffic_by_campaign_channel.plot(kind="bar", color='#FFC107',
        ↪edgecolor="black")

for bar in ax.containers:
    ax.bar_label(bar, fmt='%.3f', label_type='edge', padding=3, color='black',
        ↪fontsize=12)

plt.xlabel("Type of marketing channel")
plt.ylabel("Average number of website visits")
plt.title("Impact of advertising campaign type on website traffic")
plt.xticks(rotation=0)
plt.ylim(20, 30)
plt.show()
```



What percentage of visitors (WebsiteVisits) view more than one page (PagesPerVisit)?


```
[622]: percentage_more_than_one_page = (df[df["PagesPerVisit"] > 1]["WebsiteVisits"].
      ↪sum() / df["WebsiteVisits"].sum()) * 100
percentage_more_than_one_page
```

```
[622]: 100.0
```

How does the duration of time spent on a website (TimeOnSite) affect conversion probability?

```
[625]: converted = df[df["Conversion"] == 1]["TimeOnSite"]
not_converted = df[df["Conversion"] == 0]["TimeOnSite"]

plt.figure(figsize=(8, 6))
sns.histplot(converted, color="green", label="Conversion", kde=True, bins=30)
sns.histplot(not_converted, color="red", label="No conversion", kde=True,
      ↪bins=30)
plt.xlabel("Time on site (minutes)")
plt.ylabel("Frequency")
plt.title("Impact of time on site on conversion")
plt.legend()
plt.show()
```



Which advertising campaigns (CampaignType) contribute to longer time spent on the website?


```
[628]: time_by_campaign_type = df.groupby("CampaignType")["TimeOnSite"].mean().
        ↪sort_values(ascending=False)

plt.figure(figsize=(12, 6))
ax = time_by_campaign_type.plot(kind="bar", color='pink', edgecolor="black")

for bar in ax.containers:
    ax.bar_label(bar, fmt='%.3f', label_type='edge', padding=3, color='black',
    ↪fontsize=12)

plt.xlabel("Type of advertising campaign")
plt.ylabel("Average time on site (minutes)")
plt.title("Impact of advertising campaigns on time spent on site")

plt.xticks(rotation=0)
plt.ylim(0, 10)
plt.show()
```




Which advertising channels (CampaignChannel) contribute to longer time spent on the website?

```
[631]: time_by_campaign_channel = df.groupby("CampaignChannel")["TimeOnSite"].mean().
        ↪sort_values(ascending=False)

plt.figure(figsize=(12, 6))
ax = time_by_campaign_channel.plot(kind="bar", color='yellow',
        ↪edgecolor="black")

for bar in ax.containers:
    ax.bar_label(bar, fmt='%.3f', label_type='edge', padding=3, color='black',
        ↪fontsize=12)

plt.xlabel("Type of marketing channel")
plt.ylabel("Average time on site (minutes)")
plt.title("Impact of advertising campaigns on time spent on site")
plt.xticks(rotation=0)
plt.ylim(0, 10)
plt.show()
```




Which advertising campaigns (CampaignType) contribute to a higher number of page views?

```
[634]: pages_by_campaign_type = df.groupby("CampaignType")["PagesPerVisit"].mean().
        ↪sort_values(ascending=False)

plt.figure(figsize=(12, 6))
ax = pages_by_campaign_type.plot(kind="bar", color='lightgreen',
        ↪edgecolor="black")

for bar in ax.containers:
    ax.bar_label(bar, fmt='%.3f', label_type='edge', padding=3, color='black',
        ↪fontsize=12)

plt.xlabel("Type of advertising campaign")
plt.ylabel("Average number of page views")
plt.title("Impact of campaign type on page views")
plt.xticks(rotation=0)
plt.ylim(0, 10)
plt.show()
```


Which advertising channels (CampaignChannel) contribute to a higher number of page views?

```
[637]: pages_by_campaign_channel = df.groupby("CampaignChannel")["PagesPerVisit"].
        ↪mean().sort_values(ascending=False)

plt.figure(figsize=(12, 6))
ax = pages_by_campaign_channel.plot(kind="bar", color='blue', edgecolor="black")

for bar in ax.containers:
    ax.bar_label(bar, fmt='%.3f', label_type='edge', padding=3, color='black',
        ↪fontsize=12)


plt.xlabel("Type of marketing channel")
plt.ylabel("Average number of page views")
plt.title("Impact of marketing channels on page views")
plt.xticks(rotation=0)
plt.ylim(0, 10)
plt.show()
```



How does the number of website visits affect conversion?

```
[640]: conversion_by_visits = df.groupby("WebsiteVisits")["Conversion"].mean().
        ↪reset_index()

plt.figure(figsize=(10, 6))
sns.lineplot(x="WebsiteVisits", y="Conversion", data=conversion_by_visits,
        ↪marker="o", color="blue")
plt.xlabel("Number of website visits")
plt.ylabel("Average conversion rate")
plt.title("Impact of website visits on conversion")
plt.show()
```




Conversion funnel analysis

```
[643]: stages = {
    "Website Visits": df["WebsiteVisits"].sum(),
    "Engagement (PagesPerVisit > 1)": df[df["PagesPerVisit"] > 1][
        "WebsiteVisits"].sum(),
    "Interest (EmailOpens)": df["EmailOpens"].sum(),
    "Clicks (EmailClicks)": df["EmailClicks"].sum(),
    "Conversion": df["Conversion"].sum()
}

funnel_df = pd.DataFrame(list(stages.items()), columns=["Stage", "Users"])

plt.figure(figsize=(8, 6))
plt.barh(funnel_df["Stage"], funnel_df["Users"], color="royalblue")
plt.xlabel("Number of users")
plt.ylabel("Funnel stages")
plt.title("Conversion funnel analysis")
plt.gca().invert_yaxis()
plt.show()
```



The main issue is the low email open rate. If users don't open emails, they never reach the click stage, which impacts overall conversions. Possible reasons for low open rates: Unattractive subject line – If it doesn't grab attention, the email remains unopened. Sending time – If the email arrives at an inconvenient time, users may overlook it. Trust in the sender – The sender's address might not inspire confidence, or the email may look like spam. Email fatigue – If users receive too many emails, yours might get lost in the clutter. What can be improved: Test subject lines – Conduct A/B testing to determine which topics attract attention best. Optimize sending time – Analyze audience activity and send emails at the most convenient hours. Increase personalization – Emails that feel personal are opened more often. Use recipients' names, interaction history, and personalized recommendations. Improve email reputation – Ensure emails come from a recognizable, trustworthy address. Use intriguing content – Create a sense of value and curiosity in the first sentence. The issue isn't just about open rates—it also affects click-through rates. Enhancing the email text, call-to-action (CTA), and design can further improve engagement.

Which factors are most important for conversion?

We encode categorical variables using the One-Hot Encoding method since these variables do not have a natural order.

```
[648]: df = df.drop(columns=['AdSpendGroup', 'AgeGroup', 'IncomeGroup', 'Efficiency', 'new_customer_ratio', 'retention_ratio'])
df = pd.get_dummies(df, columns=['Gender', 'CampaignChannel', 'CampaignType'])
bool_columns = df.select_dtypes(include=['bool']).columns
df[bool_columns] = df[bool_columns].astype(int)
```

We normalize numerical variables to ensure consistency in scale. For variables with a small range, we use Min-Max Scaler. For other data, since they do not follow a normal distribution and contain

outliers, we first apply Robust Scaler and then Min-Max Scaler to bring them to a common scale.

```
[651]: scaler = RobustScaler()
df['Age'] = scaler.fit_transform(df[['Age']])
df['Income'] = scaler.fit_transform(df[['Income']])
df['AdSpend'] = scaler.fit_transform(df[['AdSpend']])
df['WebsiteVisits'] = scaler.fit_transform(df[['WebsiteVisits']])
df['SocialShares'] = scaler.fit_transform(df[['SocialShares']])
df['LoyaltyPoints'] = scaler.fit_transform(df[['LoyaltyPoints']])

scaler = MinMaxScaler()
df['PagesPerVisit'] = scaler.fit_transform(df[['PagesPerVisit']])
df['TimeOnSite'] = scaler.fit_transform(df[['TimeOnSite']])
df['EmailOpens'] = scaler.fit_transform(df[['EmailOpens']])
df['EmailClicks'] = scaler.fit_transform(df[['EmailClicks']])

df['Age'] = scaler.fit_transform(df[['Age']])
df['Income'] = scaler.fit_transform(df[['Income']])
df['AdSpend'] = scaler.fit_transform(df[['AdSpend']])
df['WebsiteVisits'] = scaler.fit_transform(df[['WebsiteVisits']])
df['SocialShares'] = scaler.fit_transform(df[['SocialShares']])
df['LoyaltyPoints'] = scaler.fit_transform(df[['LoyaltyPoints']])

[653]: X = df.drop(columns=['Conversion', 'CustomerID'])
y = df['Conversion']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

model = XGBClassifier(random_state=42)

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

feature_importance = model.feature_importances_
features = X.columns

features_sorted = [x for _, x in sorted(zip(feature_importance, features))]
importance_sorted = sorted(feature_importance)


plt.figure(figsize=(12, 10))
```

```
plt.scatter(importance_sorted, features_sorted, color='blue', s=100)
plt.xlabel("Importance Score")
plt.ylabel("Features")
plt.title("Feature Importance (Scatter Plot)")
plt.grid()
plt.show()
```

Accuracy: 0.92

Classification Report:

	precision	recall	f1-score	support
0	0.77	0.45	0.57	194
1	0.93	0.98	0.95	1406
accuracy			0.92	1600
macro avg	0.85	0.72	0.76	1600
weighted avg	0.91	0.92	0.91	1600



```
[654]: y_prob = model.predict_proba(X_test)[:, 1]
```

```

precision, recall, thresholds = precision_recall_curve(y_test, y_prob)


f1_scores = 2 * (precision * recall) / (precision + recall)
optimal_threshold = thresholds[f1_scores.argmax()]

print("Optimal threshold:", optimal_threshold)

plt.plot(recall, precision, label="Precision-Recall Curve")
plt.xlabel("Recall")
plt.ylabel("Precision")
plt.title("Precision-Recall Curve")
plt.legend()
plt.show()

```

Optimal threshold: 0.528863



```

[656]: param_grid = {
    "max_depth": [3, 5, 7, 10],
    "learning_rate": [0.01, 0.1, 0.2, 0.3],
    "n_estimators": [50, 100, 200],

```

```

        "subsample": [0.8, 1],
        "colsample_bytree": [0.8, 1]
    }

    grid_search = GridSearchCV(XGBClassifier(random_state=42), param_grid, cv=3,
        ↳scoring="accuracy", n_jobs=1)
    grid_search.fit(X_train, y_train)

    print("Best parameters:", grid_search.best_params_)

```

Best parameters: {'colsample_bytree': 1, 'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 200, 'subsample': 1}

```

[658]: X = df.drop(columns=["Conversion", 'CustomerID'])
        y = df["Conversion"]

        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
            ↳random_state=42)

        model = XGBClassifier(random_state=42, colsample_bytree=1, learning_rate=0.1,
            ↳max_depth=3, n_estimators=200,
                                subsample=1, reg_alpha=0.1, reg_lambda=1.0)
        model.fit(X_train, y_train)

        y_proba = model.predict_proba(X_test)[: , 1]

        optimal_threshold = 0.5558796
        y_pred = (y_proba >= optimal_threshold).astype(int)

        accuracy = accuracy_score(y_test, y_pred)
        print(f"Accuracy: {accuracy:.2f}")

        print("\nClassification Report:")
        print(classification_report(y_test, y_pred))

        feature_importance = model.feature_importances_
        features = X.columns

        features_sorted = [x for _, x in sorted(zip(feature_importance, features))]
        importance_sorted = sorted(feature_importance)

        plt.figure(figsize=(12, 10))
        plt.scatter(importance_sorted, features_sorted, color="blue", s=100)
        plt.xlabel("Importance Score")
        plt.ylabel("Features")
        plt.title("Feature Importance (Scatter Plot)")
        plt.grid()

```




```
plt.show()
```

Accuracy: 0.94

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.59	0.69	194
1	0.95	0.98	0.96	1406
accuracy			0.94	1600
macro avg	0.89	0.79	0.83	1600
weighted avg	0.93	0.94	0.93	1600



Thus, the most important factors for conversion are CampaignType_Conversion, EmailClicks, PreviousPurchases, TimeOnSite, and EmailOpens. To maximize results, focus on the following aspects: CampaignType_Conversion – If a specific campaign type leads to better conversions, increasing investment in scaling it can be beneficial. EmailClicks & EmailOpens – Optimizing email content (headlines, CTA, personalization) is crucial to increasing open and click rates. PreviousPurchases – Analyzing customer purchase behavior helps personalize offers and encourage repeat sales. TimeOnSite – A high metric may indicate interest, but if conversions are low, UX/UI improvements and stronger CTAs should be considered.

Budget optimization

```
[661]: df = pd.read_csv('digital_marketing_campaign_dataset.csv')

[662]: factor_importance = pd.Series(model.feature_importances_, index=X_train.columns)
factor_importance /= factor_importance.sum()

channel_performance = df.groupby("CampaignChannel").agg(
    Conversion=("Conversion", "sum"),
    ConversionRate=("ConversionRate", "mean"),
    ClickThroughRate=("ClickThroughRate", "mean"),
    AdSpend=("AdSpend", "sum")
)

campaign_performance = df.groupby("CampaignType").agg(
    Conversion=("Conversion", "sum"),
    ConversionRate=("ConversionRate", "mean"),
    ClickThroughRate=("ClickThroughRate", "mean"),
    AdSpend=("AdSpend", "sum")
)

valid_factors = ["Conversion", "ConversionRate", "ClickThroughRate"]
available_factors = [f for f in valid_factors if f in factor_importance.index]

channel_weights = sum(
    factor_importance[factor] * (channel_performance[factor] /
    ↪channel_performance[factor].sum())
    for factor in available_factors
)

campaign_weights = sum(
    factor_importance[factor] * (campaign_performance[factor] /
    ↪campaign_performance[factor].sum())
    for factor in available_factors
)

channel_weights = np.array(channel_weights)
campaign_weights = np.array(campaign_weights)

total_budget = df["AdSpend"].sum()

def optimize_budget(weights, current_spend):
    constraints = {'type': 'eq', 'fun': lambda budget_allocation: np.
    ↪sum(budget_allocation) - total_budget}
    bounds = [(0.01 * total_budget, 0.8 * total_budget)] * len(weights)
    initial_budget = current_spend.to_numpy() if isinstance(current_spend, pd.
    ↪Series) else np.array(current_spend)
```

```

    result = minimize(lambda b: -np.dot(weights, b), initial_budget,
↳ bounds=bounds, constraints=constraints, method='SLSQP')
    return result.x

optimized_channel_budget = optimize_budget(channel_weights,
↳ channel_performance["AdSpend"])
optimized_campaign_budget = optimize_budget(campaign_weights,
↳ campaign_performance["AdSpend"])

channel_budget_distribution = {channel: round(optimized_channel_budget[i], 2)
↳ for i, channel in enumerate(channel_performance.index)}
campaign_budget_distribution = {campaign: round(optimized_campaign_budget[i],
↳ 2) for i, campaign in enumerate(campaign_performance.index)}

```

```

[663]: plt.figure(figsize=(10, 5))
bars = plt.bar(channel_budget_distribution.keys(), channel_budget_distribution.
↳ values(), color='blue')

for bar in bars:
    plt.text(
        bar.get_x() + bar.get_width()/2, bar.get_height(),
        f"{bar.get_height():.2f}", ha='center', va='bottom', fontsize=10
    )

plt.xlabel("Promotion Channel")
plt.ylabel("Budget ($)")
plt.title("Optimized Budget Allocation by Channel")
plt.xticks(rotation=45)
plt.grid(axis="y", linestyle="--", alpha=0.7)
plt.show()

plt.figure(figsize=(10, 5))
bars = plt.bar(campaign_budget_distribution.keys(),
↳ campaign_budget_distribution.values(), color='green')

for bar in bars:
    plt.text(
        bar.get_x() + bar.get_width()/2, bar.get_height(),
        f"{bar.get_height():.2f}", ha='center', va='bottom', fontsize=10
    )

plt.xlabel("Campaign Type")
plt.ylabel("Budget ($)")
plt.title("Optimized Budget Allocation by Campaign Type")
plt.xticks(rotation=45)
plt.grid(axis="y", linestyle="--", alpha=0.7)
plt.show()

```

